

Comment in Response to the Copyright Office Request for Comments on Artificial Intelligence and Copyright (August 30, 2023)

Federal Register: August 30, 2023 (Volume 88, Number 167)

Notices – Page 59942-59949

LIBRARY OF CONGRESS

The United States Copyright Office

DEPARTMENT OF COMMERCE

National Telecommunications and Information Administration

This comment is filed by Brian E. Lewis, an American citizen and member of the Washington State Bar, as a representative of Scenario, Inc.

I would like to express my gratitude to the Copyright Office for affording the public the opportunity to provide input.

Contact Information:

Brian E. Lewis, Managing Partner

Azora Law, PLLC

15600 NE 8th St

Suite B1-178

Bellevue, WA 98008

(206) 910-6574

blewis@azoralaw.com

INTRODUCTION TO COPYRIGHT AND AI

Executive Summary including Recommendations

I. Introduction

The emergence of AI-generated art marks a transformative juncture with far-reaching implications for every area of the arts. Initially, these programs may displace some fundamental roles that sustain artists, and the resulting hardship for these creators should not be underestimated or overlooked. Earning a living in the art world is already a challenging feat, and the potential for AI to supplant these roles is bound to disrupt the status quo. However, this disruption parallels historical shifts, such as when painters grappled with the advent of cameras or when piano players faced the invention of player pianos. See *White-Smith Music Pub. Co. v. Apollo Co.*, 209 U.S. 1, 28 S. Ct. 319, 52 L. Ed. 655 (1908). The art world will adapt and use AI generative techniques to their advantage. Artists will not disappear, as long as we are careful, in this disruptive time, of protecting the rights of those artists in their creations. It's crucial to ensure the rights of artists in their creations are safeguarded, ensuring that artists continue to be a vital force in the creative landscape.

II. Statement of Interest

We are writing on behalf of Scenario, Inc. ("Scenario"), a provider of richly guided generative AI art services used predominantly in the game industry. Unlike many generative AI platforms, the Scenario platform requires clients to use a blend of sources. As with all generative AI platforms, one segment of data is the widely sourced foundational base model that is a statistical abstraction of potentially millions of works. Unlike many platforms, the skilled artist using the Scenario platform also draws on a small number of predominating artworks (often specifically sourced and licensed from the human artist). Importantly, the artist using the Scenario software must have a clear intention of what the final product is to look like in order to manipulate the tools in the Scenario software to produce a coherent engineered family of related derivative works.

This allows an existing character to be rapidly adapted in extremely little time to the iconic attributes and personality of that character. While this process is very efficient, it still requires significant input, intention, and engineering by the Art Director to manage the

generative process. Thus significant human input is required both in creating the predominating artwork style and features for a character by an artist, and the nuanced and skilled guidance of an Art Director skilled at using the software. To help demonstrate the concept of an artist having the skills and tools and intention to create a family of coherent derivative works, we provide the following example in which the predominating artwork is presented in the left most column and the output using well known unimodal generative AI tools are in the center, and a multimodal generative AI tool is used to create a family of related derivative works having a coherent and sufficiently delineated set of attributes to merit protection is at right:



The game industry stands at a unique juncture. Game design is a creative form of artistic expression that requires considerable skill and training to effectively master. This art of designing the rules, algorithms and heuristics for player interaction provides enormous benefit to society through teaching critical thinking, math skills, and statistics, and it also reduces suffering by promoting human connection. Never in human history has there been more creative energy

(both in absolute terms and as a percentage of the population) been more creative energy put into game design. Some of the most brilliant minds in the world are now able to dedicate their lives to benefiting society solely by practicing the art of designing new games. We are providing a sufficiently robust economic incentive to lure writers, philosophers, doctors of physics and mathematics, and others to specialize in the art and science of designing game play.

The art of game design is unique because there is currently essentially no statutory intellectual property protection for those practicing this craft. Notwithstanding this lack of legal protection, the rate of evolution in game design and the remuneration for those practicing the craft are the highest they have ever been. In short, progress of this useful art and science are arguably optimally assured in a world where the artwork being created (the design of unique and original game engines) has almost no statutory protection.

Over the past few decades, games have evolved into a burgeoning art form, merging various artistic disciplines. Video games require the intricacies of software coding, a compelling narrative, intricate visual artistry, and a rich soundscape, often culminating in a complete musical score. Currently, the landscape of the industry is such that only the most financially robust companies can afford to develop video games because they are the only ones that can afford to hire all the artists necessary. This functional oligopoly of companies will probably continue to engage artists for tailor-made works that precisely align with their creative vision. However, the integration of richly human guided generative AI as an option fundamentally changes the paradigm, opening doors for smaller developers to initiate their projects.

This shift is monumental. It democratizes the process, enabling a wider array of creative minds to contribute to the tapestry of video game artistry. Smaller studios and independent developers, who may have previously struggled to assemble extensive artistic teams, can now harness the power of generative AI to kickstart their endeavors. This means a richer, more diverse array of games can come to fruition, reflecting a broader spectrum of creative voices and visions.

Moreover, this transformation benefits the gaming community as a whole. It means a greater diversity of games, catering to a wider range of tastes and interests. It fosters innovation, as developers can experiment with new concepts and styles without the same financial constraints. Ultimately, this shift has the potential to not only empower smaller developers but

also invigorate the entire video game industry, ushering in a new era of creativity and accessibility.

The drafters of our Constitution saw value in incentivizing the creation of more art, and if we are careful with this emerging technology, that is precisely what richly human guided generative AI programs will foster. In the following comments, this memorandum will explain why the use of works for the purpose of training AI programs constitutes fair use and why imposing copyright liability on providers of AI programs, rather than the users of the AI tools, will stifle innovation and hurt artists.

III. Constitutional Basis and the Purpose of Copyright

Copyright law in the United States finds its foundation in the nation's Constitution. The intellectual property clause of the U.S. Constitution grants “Congress the power to promote the Progress of Science and the useful Arts, by securing for limited Times, to Authors and Inventors, the exclusive Right to their respective Writings and Discoveries.” U.S. Const. Art. I, § 8, Clause 8.

In other words, the intent of the drafters was to allow Congress to give a temporary monopoly to creators so that they may financially benefit from sharing their works. Without this monopoly, an artist risks unfettered copying of their works and sale at a lower price by others who had no part in the creation of those works. The drafters of the Constitution saw a benefit in more art being available to the public and saw copyright as the way to incentivize artists. All discussion of what is to be done with copyright in regard to AI must keep in mind this intent of “promot[ing] the Progress of Science and the useful Arts” by granting a limited monopoly to “Authors.”

IV. Recommendations

We recommend that the degree of creative human intention in conceptualizing and creating a finished work be a crucial factor in determining whether a work proposed for copyright registration.

Unimodal generative AI platforms create artwork based on human input in a single format. Unimodal systems often generate a work based on a few human words such as “angry bird”. We request that the work product of such systems not be granted copyright protection because the level of creative human input, design, intention, and originality is de minimis.

Richly human guided multimodal generative AI platforms such as the one offered by Scenario create artwork based on a small number of statistically predominating artworks. The model then requires the skilled guidance of an artist to efficiently expand the guidance of the predominating artworks to a limitless variety of settings, actions, styles, and interactions all with a coherent look and feel that is sufficiently delineated expression. We request that guided multimodal input based generative AI platforms have the right to be granted copyright protection on their resulting works.

To achieve the Copyright Office’s administrative need for efficient processing, we counsel that any applicant using generative AI in creation of a work be required to attest to the system used in creation of the work. Works created using unimodal generative AI platforms would be denied registration. This denial could be on a rebuttable presumption basis that allows the applicant to demonstrate significant time and energy went into creating the work sought to be protected. Conversely, works created using multimodal generative AI platforms would be allowed registration; however, such registration would be subject to judicial review in which the burden would be on the applicant to demonstrate the level of human effort should there be a challenge to the enforceability of copyright protection in the subject work.

ANSWERS TO SPECIFIC QUESTIONS

Training

Q 6: What kinds of copyright-protected training materials are used to train AI models, and how are those materials collected and curated?

Question 6.1: How or where do developers of AI models acquire the materials or datasets that their models are trained on? To what extent is training material first collected by third-party entities (such as academic researchers or private companies)?

Richly human guided multimodal generative AI platforms, such as the one created by Scenerio.com, use two distinct types of data sources and may require considerable human art direction in the creation of works. The data sources include vast generic data embodied as a statistical abstraction of all the written and visual arts evaluated by the AI learning model (“*Foundational Base Model*”). The data sources also include the statistically predominating artworks that define the physical, conceptual, and stylistic qualities that allow the designated predominating artworks to sufficiently delineate the resulting adaptation to contain unique elements of expression most viewers would associate with the original (“*Predominating Data*”). They commission works from human artists, properly obtain derivative work rights, and then train their model on those works.

Question 6.2: To what extent are copyrighted works licensed from copyright owners for use as training materials? To your knowledge, what licensing models are currently being offered and used?

Foundational Base Model that comprises part of the AI learning creates a statistical background for creation of a work. Foundational Base Model is statistically abstracted. If one was to look at the Mona Lisa, one might say that the eyes had pupils of a certain diameter, centered a given distance apart, with one being slightly higher relative to the centerline of the face than the other. These statistical descriptions are not a copy of the artwork itself, but a relational statistical model of features, colors, and other attributes. Foundational Base Model is neither copied nor licensed from the creators.

Predominating Data, data that embodies the few works that are intended to have statistically predominating influence over the resulting work, are based on commissioned works of human created art. The artist has knowingly granted and been compensated for providing the right to create derivative works based on their artwork.

This two-step approach to data sources assures that no work is scanned as a complete work of art without the artist granting rights to use the work and create derivatives thereof. Statistical abstractions of art are made without the creator's consent, but those works number in the millions, no direct copy of any one artwork is ever created, and there is no possibility that the resulting work would be substantially similar to any one element of Foundational Base Model.

Question 6.4: Are some or all training materials retained by developers of AI models after training is complete, and for what purpose(s)? Please describe any relevant storage and retention practices.

The Scenario application only retains Predominating Data, data that embodies the few works that are intended to have statistically predominating influence over the resulting work, are retained in memory to allow the Art Director to visually compare derivative works created by the platform to the original embodiments of the character or other source product. The artist has knowingly granted and been compensated for providing the right to create derivative works based on their artwork. Conversely, the millions of works of art used to train the Foundational Base Model are immediately converted into statistical abstracts and the original work of art as a coherent image is not saved anywhere.

Q 7: To the extent that it informs your views, please briefly describe your personal knowledge of the process by which AI models are trained. The Office is particularly interested in:

Question 7.1: How are training materials used and/or reproduced when training an AI model? Please include your understanding of the nature and duration of any reproduction

of works that occur during the training process, as well as your views on the extent to which these activities implicate the exclusive rights of copyright owners.

One challenge emerging from new technologies lies in determining the legal threshold for when a copy is deemed to have been created. In the digital realm, nearly all copyrightable works can exist in a digital format, prompting courts to grapple with questions about whether temporary storage in computer memory amounts to the production of a new copy.

For instance, the first-sale doctrine traditionally enables owners of a copy to freely sell it on the open market without incurring royalties. This principle hinges on the notion that no new copy is generated, and the original creator has already reaped the benefits of the initial sale. However, this doctrine doesn't extend to digital copies or phonorecords.

A notable case, *Capitol Recs., LLC v. ReDigi Inc.*, 934 F. Supp. 2d 640, S.D.N.Y. 2013, aff'd, 910 F.3d 649, 2d Cir. 2018, involved ReDigi's development of software facilitating legal consumer-to-consumer sales of digital music. Despite efforts to ensure only one copy circulated, the court determined that this process inevitably resulted in the creation of a new copy, constituting a violation of copyright law.

Similarly, the Ninth Circuit has ruled that even temporary copying of a work in a computer's memory constitutes an act of copying amounting to potential copyright infringement. *MAI Sys. Corp. v. Peak Computer, Inc.*, 991 F.2d 511, 9th Cir. 1993.

The Second Circuit has provided some insight into the duration required for a copy to be considered transitory. In *Cartoon Network LP, LLLP v. CSC Holdings, Inc.* 536 F.3d 121, 2d Cir. 2008, data buffered for no more than 1.2 seconds was deemed fleeting enough to qualify as a transitory period.

Richly human guided multimodal generative AI platforms, such as the one created by Scenerio.com, use Foundational Base Model (statistical models of millions of works on which the AI system is trained but those works are not held in memory) and Predominating Data (i.e. images intended to have a statistically predominating influence on the resulting work to produce a coherent and sufficiently delineated product identifiable as being an adapted representation of the original).

Foundational Base Model is not held in memory, but Predominating Data is licensed from the artist and held in memory for an extended period. The use of Foundational Base Model does not infringe the rights of the creator, and the use of Predominating Data would infringe the rights of the creator unless those rights are expressly licensed to the user.

Question 7.2: How are inferences gained from the training process stored or represented within an AI model?

The distributable version of our AI model does not retain any of the training data within it. After the model undergoes training on an image, a process known as "stable diffusion" is initiated. This process allows the training data to be removed from memory and replaced with a statistical abstraction generated by the program based on what it learned during training.

To put this in human terms, imagine learning how to draw faces by closely replicating faces from a magazine. Then, you're tasked with drawing "a brunette woman with green eyes." You've learned from various copyrighted works how to depict brown hair, femininity, and eyes. Over time, you might not recall the exact images you initially learned from, but you'll still possess the ability to draw faces.

Computers admittedly possess a more photographic memory than humans, so to simulate this "forgetting" of the original image, we train the system on how to draw something and then instruct it to do so. If it has been trained effectively, it will be capable of drawing what you've instructed, and the output won't resemble the original image any more than two separate individuals asked to draw a face. Both drawings will feature two eyes, a nose, and a mouth, but any other similarities arise from the choices of the artist.

It's worth noting that this underscores the importance, particularly for the owners of original copyrighted works, of training an AI model on a diverse range of samples. If you exclusively train the model to draw brunettes using images of Michelle Obama, every brunette it attempts to draw will resemble her. Similarly, if your sole reference for "science fiction" writing is "Ender's Game," a generated science fiction novel would likely encroach on Orson Scott Card's copyright in the novel. This is true even if the original work is not explicitly stored within the AI model after training is complete.

Q 8: Under what circumstances would the unauthorized use of copyrighted works to train AI models constitute fair use? Please discuss any case law you believe relevant to this question.

Even with a substantial training set, there's a colorable argument that using an image for training inherently creates a new copy. It's suggested that the latent image produced post-training is essentially a derivative work, and any subsequent interpolation of a new image is based solely on these derivative works. Both the creation of a copy and the development of these derivatives would likely encroach upon the copyrights of the original owners.

Nonetheless, the use of these images could be deemed fair use under current legal standards. The initial training of the model results in a copy, but this is employed in a profoundly transformative manner. Its application has shifted from visual art to serving as input for training an AI model. The courts have previously dealt with similar scenarios, notably cases involving Google's compilation of copyrighted works for its search engine. See *Perfect 10 v. Google, Inc.*, 416 F. Supp. 2d 828 (C.D. Cal. 2006). See also *Authors Guild v. Google, Inc.*, 804 F.3d 202 (2d Cir. 2015). In both instances, the court determined that the use of the works was profoundly transformative and provided a search function of immense public value. Moreover, this use did not significantly undermine the market for the works and endowed them with a new purpose.

Similarly, in this context, the AI model duplicates the images strictly for program training, enabling individuals without artistic training but with creative ideas to generate new works. This holds significant public value and directly aligns with the constitutional aim of the intellectual property clause – "to promote the Progress of Science and the useful Arts". In this context, copying the works serves the purpose of training the software. The knowledge extracted from these works is then utilized to generate more art, representing an entirely new and indisputably transformative use of the original works. Additionally, employing these works for machine learning doesn't impede artists' ability to manage and gain from their original creation. Thus, the use of copyrighted works solely for AI program training should always be considered fair use.

Question 8.1: In light of the Supreme Court’s recent decision in *Google v. Oracle America* and *Andy Warhol Foundation v. Goldsmith*, how should the “purpose and character” of the use of copyrighted works to train an AI model be evaluated? What is the relevant use to be analyzed? Do different stages of training, such as pre-training and fine-tuning, raise different considerations under the first fair use factor?

In light of the recent Supreme Court decisions in *Google v. Oracle America* (141 S. Ct. 1183, 209 L. Ed. 2d 311 (2021)) and *Andy Warhol Foundation v. Goldsmith* (598 U.S. 508, 143 S. Ct. 1258, 215 L. Ed. 2d 473 (2023)), the evaluation of the "purpose and character" of using copyrighted works to train an AI model should center on the training process itself. Here, the work used to train is copied in its entirety, even if briefly, for the purpose of training a computer model. This is distinctly different from the work’s original purpose as, for example, a visual art piece.

In *Google v. Oracle America*, Google was found to have a valid fair use case where they had copied parts of code that Oracle owned in copyright. Google did this for utilitarian reasons, namely, making their software easier to use for developers already familiar with Oracle’s Java. In doing so, Google helped to unlock the future creativity of developers that would have been stifled without an accessible interface. Like in Google, the use of these artworks to train the AI model is to help unlock the creative possibilities of the future users of the AI model. It's crucial to note that AI models thrive with extensive training sets, reducing the likelihood of undue impact on any individual artist within the larger dataset. Restricting the fair use of art for AI training would be detrimental to both artists and AI developers in need of substantial training data.

Contrastingly, in *Andy Warhol Foundation v. Goldsmith*, the Court ruled against the Foundation's fair use claim. The Court held that merely overlaying colors onto a copyrighted photograph, even if done by a renowned artist like Andy Warhol, did not constitute a transformative use. Some artists harbor concerns that AI programs might replicate this, making minor alterations to their works before undermining their market. However, this apprehension is not well-founded. It underscores the necessity for AI developers to have the option of a fair use defense to access the large and diverse training sets they require. Focusing solely on training the

AI model with Goldsmith's photograph of Prince, for example, would likely yield an output that lacks the transformative element necessary to avoid infringement liability.

Question 8.4: What quantity of training materials do developers of generative AI models use for training? Does the volume of material used to train an AI model affect the fair use analysis? If so, how?

The quantity of training materials utilized in developing generative AI models can vary significantly depending on the specific case. In instances where copyrighted works are used without explicit permission for training, employing a larger pool of works for training purposes should weigh in favor of fair use.

Utilizing a diverse and extensive set of training materials not only enhances the quality of the resulting output but also benefits the original copyright owners. As mentioned earlier, if an AI model is exclusively trained on images of Michelle Obama to depict brunettes, the resulting drawings are likely to closely resemble her. Similarly, if the training data for "science fiction" writing solely consists of content from "Enders Game," generated science fiction works may inadvertently infringe on Orson Scott Card's copyright in the novel. This is despite the fact that the original work is not physically stored within the AI model.

Focusing training on a singular source necessitates a greater reliance on that source to produce diverse outputs. Conversely, a diverse dataset comprising thousands of distinct works ensures that the creative process draws from a wide range of sources, resulting in outputs that do not closely resemble any single original work. This approach underscores the importance of using a comprehensive and varied set of training materials for generative AI models.

Question 8.5: Under the fourth factor of the fair use analysis, how should the effect on the potential market for or value of a copyrighted work used to train an AI model be measured? Should the inquiry be whether the outputs of the AI system incorporating the model compete with a particular copyrighted work, the body of works of the same author, or the market for that general class of works?

Under the fourth factor of the fair use analysis, the evaluation of the effect on the potential market for or value of a copyrighted work used to train an AI model should primarily focus on whether the outputs of the AI system incorporating the model directly compete with the particular copyrighted work. This means examining whether the AI-generated content could serve as a substitute for the original work in the eyes of consumers.

In most cases, a well-trained AI model, which has been exposed to a diverse range of copyrighted material, will not produce outputs that directly compete with the specific work used for training. This is because the purpose of training an AI model is to impart a broad understanding of various styles, themes, and elements found across a wide range of works, rather than merely replicating a single copyrighted piece.

This approach aligns with the principles of fair use because it acknowledges that transformative use, where the AI-generated work differs significantly from the original and does not meet the substantial similarity test, should not be hindered by copyright restrictions. If the AI-generated work is in the style of a particular artist or genre but lacks the essential elements that would make it a substantial copy, it is not a direct competitor to the original work.

On the other hand, if the inquiry were to extend to whether the AI system competes in the market for a general class of works, it could have unintended and potentially detrimental consequences. This broader scope would potentially stifle innovation and creativity in AI development, as it could effectively ban the use of the technology altogether. This is not the intent nor the purpose of copyright law, which is meant to strike a balance between protecting the rights of creators and fostering innovation and the free flow of ideas.

In conclusion, focusing the fair use analysis on whether the AI-generated outputs directly compete with the particular copyrighted work used for training is a balanced and appropriate approach. It allows for the development and application of AI technology while respecting the rights of original creators. Expanding the scope to encompass the general class of works would risk stifling progress and innovation in the field of artificial intelligence.

Generative AI Outputs: Copyrightability

Question 18: Under copyright law, are there circumstances when a human using a generative AI system should be considered the “author” of material produced by the system? If so, what factors are relevant to that determination? For example, is selecting what material an AI model is trained on and/or providing an iterative series of text commands or prompts sufficient to claim authorship of the resulting output?

Under copyright law, there are circumstances when a human using a generative AI system can be considered the "author" of material produced by the system. This determination hinges on factors such as the level of independent creation and the degree of creative input provided by the human user.

In the case of *Burrow-Giles Lithographic Co. v. Sarony* (111 U.S. 53, 4 S. Ct. 279, 28 L. Ed. 349 (1884)), the Court addressed the issue of copyright protection for photographs. Initially, there was contention that photographs were merely mechanical reproductions of existing scenes and lacked human creativity. However, the Court recognized that the composition of a photograph involves deliberate human choices and original intellectual conceptions, meeting the threshold for copyright protection.

In a similar vein, when a user interacts with a generative AI system by selecting the training material or providing iterative series of text commands or prompts, they may be contributing minimally creative elements. This input can vary in its level of creativity, depending on the specific circumstances. In some cases, the user's input may be substantial enough to constitute a form of authorship, but this determination is contingent on the extent of their creative contributions. Ultimately, the question of authorship in the context of generative AI involves a case-by-case evaluation, considering the degree of independent creative input provided by the human user.

Question 20: Is legal protection for AI-generated material desirable as a policy matter? Is legal protection for AI-generated material necessary to encourage development of generative AI technologies and systems? Does existing copyright protection for computer code that operates a generative AI system provide sufficient incentives?

Legal protection for certain AI-generated material is not only desirable as a policy matter, but an essential starting point where significant human creativity is involved. Where unimodal input is the basis for creation, there is negligible human creativity and copyright protection should be denied. Conversely, where there is richly guided multimodal input and art direction used to produce a series of derivative works of art that form part of a sufficiently delineated style and attribute cohesion, there is more than sufficient human creativity to justify and necessitate copyright protection for the derivative works.

Furthermore, it can be very difficult to distinguish between works created with AI and those without AI involvement. AI encompasses a diverse array of tools, each contributing varying degrees of creative input. Many artists have integrated some form of AI into their creative process for years, making a blanket ban potentially confusing and counterproductive within an already intricate system. These artists should not be denied protection for their art simply because some minor level of artificial intelligence was involved in the process.

Question 21: Does the Copyright Clause in the U.S. Constitution permit copyright protection for AI-generated material? Would such protection “promote the progress of science and useful arts”? If so, how?

The Copyright Clause in the U.S. Constitution grants Congress the power “to promote the Progress of Science and the useful Arts, by securing for limited Times, to Authors and Inventors, the exclusive Right to their respective Writings and Discoveries.”

In *Burrow-Giles Lithographic Co. v. Sarony*, the Court deliberated on whether Congress could extend copyright protection to photographs, addressing the argument that photographers merely mechanically captured existing scenes without significant human creativity. The Court astutely recognized that the composition of a photograph necessitated human choices and creativity, thereby meeting the threshold for copyright protection.

Similar arguments are raised concerning copyright protection for AI-generated material. Critics question whether the human user of the AI can be considered the true author. We assert that there is no bright line rule. Drawing a parallel to the *Burrow-Giles* case, the answer remains affirmative. The Court's assertion that the Constitution is sufficiently broad to encompass

copyrights for photographs, as long as they embody original intellectual conceptions of the author, resonates in this context.

Much like with photographs, the determination of whether a work possesses enough human input for copyright protection will be assessed on a case-by-case basis. Not every AI-generated work will meet the criteria for protection, but the mere use of an AI model as a tool should not be the sole decisive factor.

Protection of appropriate AI-generated material would promote the progress of science and the useful arts because it would incentivize utilizing AI to produce more copyrightable works, providing society with more art. Additionally, protecting AI-generated works would incentivize computer scientists to continue developing the technologies behind AI.

Generative AI Outputs: Labeling or Identification

Q 28: Should the law require AI-generated material to be labeled or otherwise publicly identified as being generated by AI? If so, in what context should the requirement apply and how should it work?

Question 28.1: Who should be responsible for identifying a work as AI-generated?

The challenge with this question is that it assumes a paradigm that will not apply in many cases. Certainly with a unimodal input based generative AI engine requiring only a few verbal guides, we would think that the resulting work is “AI-generated” however the distinction is far more nuanced. Artistic tools using AI systems will soon permeate our lives from images captured with our phones to Photoshop tools employed by nearly all professional photographers. To create a bright line rule that any use of generative AI tools renders the resulting work unprotectable is misguided. Given the diverse array of AI models, each with varying degrees of reliance on existing works and artistic guidance, the responsibility of identifying a work as using generative AI tools should fall on the user of the applicant seeking copyright protection. This could be effectively achieved by incorporating a prompt in the copyright registration form, where the applicant would be required to specify the primary software tools employed in creating the work. This would facilitate a swift administrative assessment by the Copyright Examiner to ascertain

the level of human input involved. While there might be some initial investigation into how different models operate, once this knowledge is established, it can be applied to subsequent cases involving the same model.

For instance, upon encountering a reference to "Scenario" an examiner could refer back to the prior investigation, and know that this particular model is heavily guided by artists and requires significant human input to produce a useful product, thereby minimizing the risk of copyright infringement. Conversely, if another model is flagged for less rigorous business practices, it could serve as an alert for potential further scrutiny when reviewing works generated using that model.

Furthermore, having users indicate in their copyright application that they utilized AI would establish a record for enforcement purposes. In cases where a copyright holder attempts to enforce their copyright and it is subsequently determined that the work lacked sufficient human input, this record could be crucial in potentially invalidating the copyright protection. This would entail a secondary review by the courts in an infringement or related proceeding to determine if the registrant should be allowed to enforce their copyright, particularly if it is found to lack the necessary human creative contribution.

This approach does not necessitate a significant overhaul of existing copyright law. If there is a substantial human creative element, the work is eligible for copyright protection. By adopting this framework, it alleviates the need for a comprehensive fair use analysis for every AI system, as the focus would be primarily on evaluating the extent of human input.