KERNOCHAN CENTER FOR LAW, MEDIA AND THE ARTS
COLUMBIA LAW SCHOOL

BEFORE THE
U.S. COPYRIGHT OFFICE

Artificial Intelligence and Copyright	**Docket No. 2023-6**

Comments of the Kernochan Center for Law, Media and the Arts

Columbia Law School

The Kernochan Center for Law, Media and the Arts appreciates the opportunity to respond to the Notice of Inquiry (NOI) concerning Artificial Intelligence and Copyright, published by the U.S. Copyright Office, 88 Fed. Reg. 167 (Aug. 30, 2023).

The Kernochan Center for Law, Media, and the Arts at Columbia Law School is one of the leading centers for intellectual property research in the United States. Its faculty and staff dedicate their research and writing to copyright, trademarks, and related areas as they concern traditional and emerging media, entertainment and the arts. The Center offers students an in-depth program of instruction, lectures, internships and externships while providing symposia, lectures, research studies and publications to the broader legal community. Founded as the Center for Law and the Arts, it was renamed in 1999 to honor Professor John M. Kernochan, its founder and a pioneer in teaching copyright in American law schools.

The Kernochan Center will not be responding to all of the questions in the NOI. Some are beyond our competence or require knowledge we do not currently possess. As to the latter, other responders' comments to this NOI may supply information sufficient to allow us to address those questions in Reply Comments. Finally, the Center is currently engaged in a study of AI inputs that overlaps with the matters raised in NOI Questions 6-17. In order not to prejudge our inquiry, we will not be responding at this time to those questions.

NOI General Questions [We reproduce the questions to which we are responding, followed by our responses]

3. Please identify any papers or studies that you believe are relevant to this Notice. These may address, for example, the economic effects of generative AI on the creative industries or how different licensing regimes do or could operate to remunerate copyright owners and/or

creators for the use of their works in training AI models. The Office requests that commenters provide a hyperlink to the identified papers.

- **AI and Authors**

Jane C. Ginsburg & Luke A. Budiardjo, Authors and Machines, 34 BERKELEY TECH. L. J. 343 (2019) *available at*: https://scholarship.law.columbia.edu/faculty_scholarship/2323

Shyamkrishna Balganesh, Causing Copyright, 117 COLUM. L. REV. 1 (2017) *available at* https://columbialawreview.org/wp-content/uploads/2017/01/01.pdf

- **Economic effects of generative AI on the creative industries**

Henry H. Perritt, Jr., "Copyright for Robots?" 57 IND. L. REV. 139 (2023) *available at* https://plus.lexis.com/api/permalink/d90da62e-1bcd-4f5c-9c37-e2c217493301/?context=1530671

Atilla Kasap, Copyright and Creative Artificial Intelligence (Ai) Systems: A Twenty-First Century Approach to Authorship of Ai-Generated Works in the United States, 19 WAKE FOREST J. BUS. & INTELL. PROP. L.REV. 335 (2019) *available at* https://www.westlaw.com/Document/I7e0951b44c6111eaadfea82903531a62/View/FullText.html?transitionType=Default&contextData=(sc.Default)&VR=3.0&RS=cblt1.0

Patrick Zurth, "Artificial Creativity? A Case Against Copyright Protection for Ai-Generated Works," UCLA J.L. & TECH., Spring 2021 *available at* https://www.westlaw.com/Document/Ibabb199f839d11ebbea4f0dc9fb69570/View/FullText.html?transitionType=Default&contextData=(sc.Default)&VR=3.0&RS=cblt1.0

- **How different licensing regimes do or could operate to remunerate copyright owners and/or creators for the use of their works in training AI models**

Mauritz Kop, Ai & Intellectual Property: Towards an Articulated Public Domain, 28 TEX. INTELL. PROP. L.J. 297 (2020). https://www.westlaw.com/Document/I931d883c5c7f11ebbea4f0dc9fb69570/View/FullText.html?transitionType=Default&contextData=(sc.Default)&VR=3.0&RS=cblt1.0

Deepak Somaya & Lav R. Varshney, "Ownership Dilemmas in an Age of Creative Machines," 36 ISSUES IN SCIENCE AND TECHNOLOGY, 79, 85, 2020. https://www.jstor.org/stable/26949112

Jenny Quang, "Does Training AI Violate Copyright Law?", 36 BERKELEY TECH. L.J. 1407 (2021). https://www.westlaw.com/Document/Ia8e44659a77811ed8636e1a02dc72ff6/View/FullText.html?transitionType=Default&contextData=(sc.Default)&VR=3.0&RS=cblt1.0

Justin Hughes, "Rules, Standards, and Copyright Fair Use" (May 4, 2020) Loyola Law School, Los Angeles Legal Studies Research Paper No. 2020-08, *available at* SSRN: https://ssrn.com/abstract=3592312

Justin Hughes, "Restating Copyright Law's Originality Requirement," 44 COLUM. J.LAW & ARTS 383 (2021), *available at* SSRN: https://ssrn.com/abstract=3811132 or http://dx.doi.org/10.2139/ssrn.3811132

Pamela Samuelson, "Generative AI meets copyright," SCIENCE 381,158-161 (2023). DOI:10.1126/science.adi0656 https://www.science.org/doi/10.1126/science.adi0656

Pamela Samuelson on The Lawfare Podcast *available at* https://www.lawfaremedia.org/article/pam-samuelson-on-copyright%27s-threat-to-generative-ai

4. Are there any statutory or regulatory approaches that have been adopted or are under consideration in other countries that relate to copyright and AI that should be considered or avoided in the United States? How important a factor is international consistency in this area across borders?

The EU has, in its 2019 Digital Single Market directive, established two exceptions for "text and datamining"; these might bear on AI inputs.[1] The first exception, in DSM Directive section 3, focuses narrowly on copying for scientific research: it covers "reproductions and extractions made by research organizations and cultural heritage institutions in order to carry out, for the purposes of scientific research, text and data mining of works or other subject matter to which they have lawful access." Section 4 sweeps more broadly and might offer a model for legally authorized constitution of training data, but the scope of the exception is unclear, and its provision for opting out may violate the Berne Convention.

Art. 4 covers:

> 1. . . . reproductions and extractions of lawfully accessible works and other subject matter for the purposes of text and data mining.
>
> 2. Reproductions and extractions made pursuant to paragraph 1 may be retained for as long as is necessary for the purposes of text and data mining.
>
> 3. The exception or limitation provided for in paragraph 1 shall apply on condition that the use of works and other subject matter referred to in that paragraph has not been

[1] The European Parliament and Council Directive on Copyright and Related Rights in the Digital Single Market 2019/790, 2019 O.J. (L. 130) 1.

expressly reserved by their rightholders in an appropriate manner, such as machine-readable means in the case of content made publicly available online.

Art. 4 imposes exceptions on "reproductions and extractions;" it does not explicitly extend to "communication to the public." As a result, the exception may not apply to AI outputs.[2] Whether or not it does, however, rightsowners may opt out of ("expressly reserve") the exception (though the means of doing so are still being developed). The opt-out poses problems of compatibility with Berne Convention norms, because it conditions the scope of protection on a declaratory obligation. Berne art 5(1) provides "The enjoyment and the exercise of these rights [including reproduction and communication to the public] shall not be subject to any formality." As detailed in Jane C. Ginsburg, *Berne-Forbidden Formalities and Mass Digitization*, 96 BOSTON U. L. REV. 745, 758-68 (2016), *available at* http://ssrn.com/abstract=2772176, a back-door limitation on the scope of protection through a formality-freighted opt-out from an exception offends Berne just as much as a front-door imposition of a formality as a condition of protection.

5. Is new legislation warranted to address copyright or related issues with generative AI? If so, what should it entail? Specific proposals and legislative text are not necessary, but the Office welcomes any proposals or text for review.

At this point, we believe it is premature to introduce new copyright legislation, but note that others' responses to this NOI may lead us to revisit that conclusion in Reply Comments. By contrast, federal legislation to protect the identity rights of performers may already be warranted.

Copyrightability

18. Under copyright law, are there circumstances when a human using a generative AI system should be considered the "author" of material produced by the system? If so, what factors are relevant to that determination? For example, is selecting what material an AI model is trained on and/or providing an iterative series of text commands or prompts sufficient to claim authorship of the resulting output?

We support the United States Copyright Office's (USCO, "the Copyright Office," or "the Office") reading of applicable precedent to suggest that a work is copyrightable only if human-authored. We also support the USCO's policy issued in March, 2023 which states that whether or not an output has a human author requires case-by-case analysis. Determining whether there is sufficient human contribution in an AI-assisted or generated output is a difficult exercise, given that human intervention can come at any stage in the process, including after the generation of an initial version of the output.

[2] Raquel Xalabader, Professor, Universitat Oberta de Catalunya, 35th Annual Horace S. Manges Lecture at Columbia Law School (Apr. 17, 2023) *available at* https://kernochan.law.columbia.edu/content/manges-lectures.

As we understand the process of production of AI-assisted/generated works, humans may intervene: (i) upstream in assembling the training data (including by creating or curating the datasets), (ii) in designing the "model" that works with the training data to generate outputs, (iii) downstream in creating prompts to elicit initial outputs and to refine them, and (iv) in selecting among or modifying outputs delivered by the system. Finally, (v) it is also possible for downstream users to supply the system with user-created or user-selected content, which the system will "treat" by executing prompts to generate new content based on the user-furnished inputs.

We understand, consistently with the approach so far articulated by the Copyright Office, that conception and control over the outputs are essential to authorship. The Copyright Office has rejected applications to register outputs whose generation was too random and uncontrolled by the prompt-supplying user. While some element of randomness does not eliminate authorship, the putative author must be able to constrain or channel the program's processing of the source material.

We believe the clearest case for authorship arises when the same human person (or group of persons together) both create or curate the training data *and* the model, and design the prompts to which the model will respond. This example does not exclude all randomness in the processing, but the definition of the universe with which the prompts will interact may sufficiently constrain the possibilities to make authorship attribution possible. (Indeed, like the early videogames cases,[3] there may be situations in which the upstream determination of the content and of the ways the model will manipulate the content so constrain the possible outcomes that the downstream user's contribution plays too slight a causal role in the generation of the outputs to attribute authorship to the end-user, and the sole authors therefore are the upstream contributors. One may think of this scenario as an extremely high-tech variation on the "Choose Your Own Adventure" series of children's books.)

Turning to **downstream authorship**, we agree with the Copyright Office that simply providing initial prompts and selecting among the choices the program offers in response does not amount to human intervention sufficient to play a proximately causal role in the generation of the output. Too much is left to chance: the user does not know what sources will inform the response, nor does she know or sufficiently control how the program will interact with the sources. The Office's analogy to *Kelley v Chicago Park District*'s characterization of the color patterns of wildflowers as too randomly generated to constitute a work of authorship seems apt: in the court's view, Mother Nature, not Mr. Kelley, provided the inputs (the seeds) and determined the outputs (the resulting flowers' wind-borne arrangement into color patterns).[4]

We believe, however, that more information about the content-generation process may be required in order to determine whether, at some point, repeatedly refined prompts may sufficiently channel the randomness of the outputs to attribute authorship of the result to the downstream user. Similarly, if the prompts edit a traditional work created or designated by the user, or if they revise a specific AI-generated output (rather than each time returning to request a

[3] See, e.g., *Williams Elecs., Inc. v. Artic Int'l, Inc.,* 685 F.2d 870 (3d Cir. 1982)
[4] *Kelley v. Chi. Park Dist*., 635 F.3d 290, 304 (2011).

new output "from scratch"), the specified starting point constrains the "autonomy" of the outputs. Does the combination of a specific starting point together with repeated and increasingly refined prompts therefore present a more persuasive case of human intervention than applying the prompts to an unknown starting point?

On the other hand, starting with an initial machine-generated output, the addition of human-controlled modifications may result in a copyrightable derivative work. For example, the end-user might edit AI-generated text, or hand-draw or use an image-editing program to alter an AI-generated image: the originality of those modifications should be assessed in the same way as are editorial or other changes to a pre-existing work.

Finally, one may raise the **question whether machine-generated outputs prompted by human users could be considered copyrightable joint works whose authors would be the prompt-supplying user and the designer(s) of the model and/or the training data**.

Alternatively, might the prompting user be considered an author on the ground that, had the upstream designers' contribution been copyrightable, then the user's contribution of prompts sufficiently elaborated to constitute a copyrightable contribution, together with her intent to merge her prompts with the operation of the AI program, would make her a co-author?

We do not believe that either of these scenarios results in a co-authorship, under the terms of the statute and its judicial interpretation.

Joint authorship: First regarding the interaction of the prompting end-user with the AI program, the asynchronous nature of the human participants' contributions precludes characterizing the outputs as a collaborative inseparable joint work.[5] As explained in detail in *Authors and Machines, supra (Question 3)*[6], the statute assumes that a merger of inseparable contributions requires the joint authors to work together. Notably, the "intent to merge" must be simultaneous with the creation of each contribution. That outcome is inherent in any actual collaboration, but as to asychronous and unacquainted contributors, it implies the kind of inchoate intent to merge (for example, the author intends that someday someone will come along and set her words to music, or vice versa) that Congress rejected for joint works in the 1976 Act.[7] When downstream users employ off-the-shelf generative AI programs, the upstream authors of the model have completed their work long before the downstream user (whom the upstream participants likely will not know and whose prompts they cannot anticipate) will have entered her prompts.

The asynchronous nature of the participation creates another impediment to considering the model designers the co-authors with the prompt-supplier: at the time of the creation of the model, the designers do not know how the program will respond to an unknown prompt, it therefore

[5] A true contemporaneous collaboration between model designers and prompt-suppliers can yield a joint work under current standards, see, e.g., *ING Presents: The Next Rembrandt*, https://perma.cc/7RS4-RM6V (visited 28 October 2023).

[6] 34 BERKELEY TECH. L.J. at 378-92

[7] See H.R. Rep. No. 94-1476 at 120 (rejecting *Shapiro, Bernstein & Co. v. Jerry Vogel Music Co*., 221 F.2d 569, 569–70 (2d Cir.1955), "Twelfth Street Rag" precedent and suggesting that such situations be treated as derivative works instead).

seems extremely attenuated to call the output the "contribution" of the upstream designers. Moreover, if, as the Second Circuit and other courts have held,[8] joint work status requires the participants to have made *copyrightable* contributions (even to inseparable works), then it becomes even more difficult to characterize randomly-generated images or text (or musical elements or code) as copyrightable. Even if the accumulated prompts of the end-user achieve a level of contribution that might constitute authorship, the results of the model designers' participation continue, from their perspective, to be equally unforeseeable (absent actual collaboration).

The practical difficulties of joint ownership also counsel against characterizing the model designers and the prompt-supplier as joint authors. Each joint author would have an undivided interest in the output, which would allow either to license on a non-exclusive basis, but would preclude either from granting exclusive rights without the other's accord. Moreover, any transfer of exclusive rights must be in writing from both joint authors, further complicating the exploitation of any copyright in the output. The Terms and Conditions (T&C) of access to the AI program might endeavor to allocate full ownership to the model designers or to the prompt-supplier, but it seems unlikely that such T&C would satisfy the statute's requirements for the transfer of copyright ownership interests delineated in §204(a).

Joint authorship minus one: We turn next to the suggestion that the prompt-supplier who provides prompts sufficiently elaborated to be considered a copyrightable contribution in their own right as a literary work becomes a co-author of the combination of her prompts with the output of the machine, on the theory that if a human had provided the responses to the prompts (intending the responses to merge with the prompts), the result would have been a joint work.

We believe this theory to have multiple shortcomings: (i) it disregards the requirement that *both* or *all* contributors to the joint work must be "authors" and thus comply with the human authorship requirement that the USCO has correctly insisted on, (ii) it overlooks the centrality of a contemporaneous "intent" to a joint work, which is unlikely to be satisfied in the absence of a non-human contributor, and (iii) it risks expanding the scope of rights in a work well-beyond what the creator of a sole-authored work might be able to obtain.

Each putative joint author, under the established rule for protectable joint works, must satisfy copyright law's "authorship" requirement. The authorship requirement has been rightly understood to require a sufficiently robust "human" element. Under the half joint author theory, this requirement would be altogether disregarded, in the process rendering it of questionable utility elsewhere as well. If a joint author does not need to be human, why insist on any author needing to be human? The only way to do so would be to accept that a contributor to a joint work need not be an author at all; which is contrary to copyright's joint works doctrine.

Joint authorship requires a contemporaneous intent, which the half joint authorship theory cannot satisfy. There is some doctrinal support for the proposition that "initial conceptualization and ongoing direction of the realization" of a work of visual art might make the supervising conceptualizer a co-author, *see CCNV v. Reid*, 846 F.2d 1485, 1497 (D.C. Cir. 1988*) aff'd. on*

[8] See, e.g, *Childress v. Taylor*, 945 F.2d 500, 506-07 (2d Cir. 1991).

other grounds, 490 U.S. 730 (1989), *citing Strauss v. Hearst Corp.,* 1988 U.S. Dist. LEXIS 1427 (S.D.N.Y. 1988).[9] But *CCNV v. Reid* concerned a *contemporaneous* collaboration in which the sculptor and the commissioning party intended their contributions to merge into a unitary whole. Moreover, citing *Strauss v. Hearst Corp*, the Second Circuit has underscored the need for "objective indices of co-authorship intent," *Thomson v. Larson* at 147 F.3d 201 n. 17, indices necessarily absent in asynchronous AI output scenarios.[10] In addition, whether the prompt-supplier "*direct[s]* the realization" of the output requires understanding the extent to which the prompts in fact "direct" the output. We currently lack that understanding. If each prompt newly rolls the dice, it is difficult to discern the dominance of will that "direction" implies, and thus hard to classify it as meeting the requirement of an objective "intent".

Treating the output as half a joint work considerably expands the scope of protection of the literary work that the aggregate of elaborated prompts might constitute. To be seen as extending to a visualization, a literary description would need to evince an extremely high degree of precision (for example, Sol LeWitt's instructions for the creation of his monumental murals, although even these included visual diagrams[11]). In effect, the literary description becomes a mode of fixation of a specific, but yet inchoate pictorial work. By contrast, a set of instructions for illustrating a literary work would likely remain ineligible for copyright protection as being an "method of operation" under §102(b).

Baker v. Selden (101 U.S. 99 (1879)) (codified in section 102(b)) and progeny[12] support the conclusion that the copyright in a written set of instructions does not confer exclusive rights to carry out those instructions. Admittedly, the instructions in that case concerned the implementation of a functional system of bookkeeping, not the creation or performance of a work of authorship, for which the instructions may be copyrightable. But the protection of architectural blueprints, stage directions, and musical scores all correspond to the creation or performance of a specific work (as do Sol LeWitt's instructions). By contrast, even highly elaborated prompts will, to our knowledge, yield multiple outputs (not all of them fully or

[9] *See also*, id: "Following CCNV's original conception of the sculpture, Snyder and other CCNV members, the district court found, monitored the progress of the work, not simply to approve Reid's embodiment of their idea, but to guide his expression and coordinate with his effort CCNV's construction of the steam grate pedestal."

[10] The impetus for the Second Circuit's high bar to finding joint author relationships was to protect the dominant author against interloping helpers: "Care must be taken . . . to guard against the risk that a sole author is denied exclusive authorship status simply because another person renders some form of assistance." Thomson v Larson 147 F.3d at 202, *quoting Childress v Taylor*, 945 F.2d 500, 504 (2d Cir. 1991).

[11] See *Sol LeWitt: A Wall Drawing Retrospective* (comprising 105 of LeWitt's large-scale wall drawings, spanning the artist's career from 1969 to 2007), https://massmoca.org/sol-lewitt/

[12] *E.g. Bikram Yoga Coll. of India, L.P. v. Evolution Yoga LLC*, 803 F.3d. 1032 (2015).

accurately responsive to the prompts).[13] Extending the scope of copyright protection in the written prompts to cover the multiplicity of potential outputs comes uncomfortably close to conferring a copyright in a method of generating images (or other works).

Where the prompt-supplier's copyright would not reach the implementation of the prompts, the half a joint work approach would in effect give the prompt author exclusive rights in a particular execution of the prompts.[14] As we understand it, the approach would not go so far as to confer rights in all the outputs that might respond to the elaborated prompts; it seems the prompt-supplier's final selection among the possibilities the AI program proffers would constitute the crowning act of (co)authorship. But we are uncertain whether the half a joint work approach can be limited to a single output. If the copyrightable contribution is the collection of elaborated prompts, with which the model designers are deemed to intend to merge whatever their model generates, then the post hoc selection among the offered options falls outside the "joint" creative process. But selection of a single output is not itself a creative act. It is not clear that a coherent stopping place exists along the continuum of response-generation. We therefore hesitate to conclude that a result impermissible through the front door of the scope of copyright can be rendered permissible through the back door of an inventive (or strained) interpretation of the authorship of a joint work.

19. Are any revisions to the Copyright Act necessary to clarify the human authorship requirement or to provide additional standards to determine when content including AI-generated material is subject to copyright protection?

At this time, we do not believe that the Copyright Act requires revision. The statute's open-ended, undefined, term "original work of authorship," is sufficiently flexible to permit dynamic interpretation; greater precision risks freezing the statute at a given point in the evolution – technological, aesthetic and otherwise – of how works are created. The

---

[13] For example, (in keeping with the season), we asked DALL-E2 to generate an image corresponding to:

a feminist witch riding a lunar-powered broomstick and wearing a green and pink cape and a yellow cap with eagles' wings, and carrying a handbag in the shape of a jack o'lantern, and a navy blue designer briefcase, and bearing on her back a cauldron containing three purple lemurs

(None of the four responses fully – or even closely – tracked the elaborate prompt)

We then tried (with equally nonresponsive responses)

three feminist witches riding a lunar-powered broomstick, one witch wearing a green and pink cape, one witch wearing a yellow cap with eagles' wings, and another witch carrying a handbag in the shape of a jack o'lantern, and all three witches bearing on their backs a cauldron containing three purple lemurs

[14] There is another way in which the half a joint work approach gives the constructive joint author more rights than under normal copyright rules: Deeming the prompter a joint author with a non-author in effect makes the end user the sole author, with full copyright ownership rights in the output.

Compendium can offer illustrative examples and guidance without locking in a particular understanding of human creativity. We believe that in general, judicial elaboration continues to be well-suited to developing a legal understanding of authorship, as courts' evolving evaluation of the copyrightability of photographs has demonstrated.

20. Is legal protection for AI-generated material desirable as a policy matter? Is legal protection for AI-generated material necessary to encourage development of generative AI technologies and systems? Does existing copyright protection for computer code that operates a generative AI system provide sufficient incentives?

We assume that by "AI-generated material" this question is inquiring about the desirability of legal protection for outputs that are not "original works of authorship." We lack the empirical information to determine whether protection for authorless outputs is "necessary to encourage development of generative AI technologies and systems," or whether "existing copyright protection for computer code that operates a generative AI system," and potentially for the compilations of training data, suffice. Nonetheless, even without that utilitarian information, normative considerations may counsel against copyright protection for authorless outputs. The courts and the Copyright Office have understood copyright to be rooted in human creativity. The constitutional term "Authors" assumes human creators; to validate the constitutionality of the 1865 amendment adding photographs to the subject matter of copyright, the Supreme Court in *Burrow-Giles v. Sarony* necessarily considered photographs to be works of human creativity rather than mere machine-generated outputs.[15] Machine-aided, to be sure, but still "master-minded" by human beings. The statute, too, contains several provisions grounded in human authorship, including the affirmation of copyright's "subsist[ence]" in original works of authorship. The acts of creation and fixation give rise to copyright. While copyright's exclusive rights may "promote the progress of science" by providing incentives to production, incentives alone do not justify exclusive rights. Were that the case, the Supreme Court should have decided *Feist v Rural Telephone* differently: the White Pages phone books were highly useful, and might on that account have warranted legal protection, but they lacked the "creative spark" of originality, and therefore did not qualify for statutory (and, the Court implied, constitutional) protection.[16]

20.1. If you believe protection is desirable, should it be a form of copyright or a separate sui generis right? If the latter, in what respects should protection for AI-generated material differ from copyright?

While we do not believe that protection of authorless outputs is consistent with the fundamentally humanist nature of copyright, it is possible that an incentive rationale (if empirically supported) could justify some form of *sui generis* protection, but it is premature to endeavor to articulate its contours. The case, empirical and theoretical, for *sui generis* protection has yet to be made.

[15] 111 U.S. 53 (1884).
[16] *Feist Publ'ns, Inc. v. Rural Tel. Serv. Co*., 499 U.S. 340, 359 (1991).

21. Does the Copyright Clause in the U.S. Constitution permit copyright protection for AI-generated material? Would such protection "promote the progress of science and useful arts"? If so, how?

As detailed in our response to Question 20, we believe the Copyright Clause implies human authorship, and that the possibility that protection might promote progress does not, standing alone, suffice to justify copyright protection.

22. Can AI-generated outputs implicate the exclusive rights of preexisting copyrighted works, such as the right of reproduction or the derivative work right? If so, in what circumstances?

We assume that this question applies not only to outputs that are fully AI-generated, but also AI-assisted. If the training data contains preexisting copyrighted works, the appearance of recognizable parts of the source works in the outputs (this is a matter in factual dispute[17]) may violate the source works' reproduction and/or derivative works rights. In any event, the manner in which the model uses the work (e.g., as a copy, as the basis for a derivative) will determine which rights are triggered; so too will the very category of work being used by the model (thus, the use of a photograph may trigger the public display right while the use of a movie clip the public performance right). Without that information, this discussion is necessarily abstract and assumes an unrealistic degree of homogeneity among AI models and their use of training data. Similarly. whether the incorporation of portions of the source works in the outputs is a fair use cannot be answered in the abstract.

Prompting an AI system to write a new story incorporating a well-delineated character (e.g., Harry Potter) may infringe the derivative works right in the character (to the extent literary characters enjoy an independent copyright), or in the literary works in which the character appears, even if the new story otherwise consists of "new matter." Similarly, prompting the system to produce images of a visual character (e.g. Spiderman) will likely (absent fair use) violate the reproduction and derivative works rights in the pictorial, graphic or sculptural work, even if the setting in which the visual character appears does not copy from preexisting works. (In the case of visual characters, their unauthorized reuse in AI-generated outputs could also pose the same trademarks issues as arise with recycling characters in analog environments.)

23. Is the substantial similarity test adequate to address claims of infringement based on outputs from a generative AI system, or is some other standard appropriate or necessary?

At this point, we lack the information that could enable us to determine whether a different standard of substantial similarity, or different means of proving the current standard, are warranted. It is possible that the "ordinary lay observer" standard does not fit, or, relatedly, that the role of expert testimony not be confined to proof of copying, but should carry over to the

[17] See, e.g., *Kadrey v. Meta Platforms, Inc.* Case 3:23-cv-03417 (N.D. Ca., filed July 7, 2023); *Getty Images v. Stability AI, Inc.*, Case 3:23-cv-00201 (D. Del. filed Feb. 3, 2023); *Anderson v. Stability AI Ltd.*, Case No. 23-cv-00201-WHO (N.D. Cal., October 30, 2023) (granting motion to dismiss but giving plaintiffs "leave to amend to provide clarity regarding their theories of how each defendant separately violated their copyrights, removed or altered their copyright management information, or violated their rights of publicity and plausible facts in support.").

issue of illicit appropriation. It is also possible that a showing of access and probative similarity could give rise to a presumption of illicit appropriation. In any event, we would need a better understanding of the inadequacies – if any – of the current test and its modes of proof.

24. How can copyright owners prove the element of copying (such as by demonstrating access to a copyrighted work) if the developer of the AI model does not maintain or make available records of what training material it used? Are existing civil discovery rules sufficient to address this situation?

We assume this question addresses proof of copying relative to the AI system's outputs. Where the outputs reveal probative similarities such as source works' watermarks[18] or other identifying indicia, such as artists' signatures,[19] copying may be inferred using the rules of circumstantial evidence.[20] If the source files contain "digital fingerprints" or copyright management information (CMI) that the outputs reproduce, these too could be considered probative similarities. But the task of identifying source works (if the developer does not maintain records) will become more difficult if the developer strips out identifying information. On the other hand, systematic removal of CMI may be actionable under section 1202(b), on the grounds that the person engaging in metadata-stripping "ha[s] reasonable grounds to know, that it will . . . conceal [the developer's own] infringement of any right under this title."[21] One court has recently rejected the contention that section 1202(b) does not reach "passive non-inclusion" of a copied file's CMI:

> Defendants argue that the complaint merely alleges "the passive non-inclusion of CMI" by neutral technology which excerpts code without the accompanying CMI, rather than the active removal of CMI from licensed code. This semantic distinction is not meaningful. Plaintiffs allege that the relevant CMI was affixed to their licensed code and that Defendants were aware that such CMI appeared repeatedly across the data used to train Codex and Copilot. . . . Defendants subsequently trained these programs to ignore or remove CMI and therefore stop reproducing it. . . . Defendants knew that these programs reproduced training data as output.[22]

The Copyright Office might facilitate proof of copying through evidence of CMI by ascertaining if the means used by rightsowners to identify their works and their authors correspond to the

---

[18] See *Getty Images (US), Inc., v. Stability AI Ltd*. (Filed Jan. 16, 2023) Case No. IL-2023-000007 (UK); *Getty Images (US), Inc. v. Stability AI, Inc.*, Case No. Case 1:23-cv-00135-GBW (D. Del.) (Filed Feb. 3, 2023).

[19] See, e.g., Shanti Escalante-De Mattei, *Artists Voice Concerns Over The Signatures In Viral LensaAI Portraits*, ARTNEWS (Dec. 9, 2022), https://www .artnews.com/art-news/news/signatures-lensa-ai-portraits-1234649633/

[20] *See* Shyamkrishna Balganesh and Peter S. Menell, *Proving Copying*, 64 WM. & MARY L. REV. 299 (2022), *available at* https://scholarship.law.wm.edu/wmlr/vol64/iss2/2

[21] A lawsuit filed against AI producer Github for systematically removing authorship attribution information, and thereby failing to abide by the attribution condition term of the General Public (free software) License, has survived a motion to dismiss, see *Doe 1 v. Github*, No. 22-CV-06823-JST, 2023 U.S. Dist. LEXIS 86983 (N.D. Cal., May 11, 2023). A defendant's concealment of its own infringement may meet the second knowledge prong of 1202(b), see *Mango v. Buzzfeed, Inc.*, 970 F.3d 167 (2d Cir. 2020).

[22] *Doe 1 v. Github*, *supra*.

categories of CMI set out in sec 1202(c). In the event that some works- or author-identifying practices do not come within sec 1202(c)(1)-(7), the Office could consider exercising its prerogative under sec 1202(c)(8) ("Such other information as the Register of Copyrights may prescribe by regulation") to list categories that fill in any gap between works- or author-identifying practices and statutorily protected CMI. In general, the Office might also encourage authors and rightsowners to include CMI in publicly disseminated copies of their works. The Copyright Office might also encourage the interested parties to develop a robust set of best practices for proof of copying.

Regarding the efficacy of civil discovery rules, unless AI platforms keep information regarding the copying of works into training data on record, there may be nothing to discover.

The European Union's pending AI Act, COM (2021) 206 final: Proposal for a REGULATION OF THE EUROPEAN PARLIAMENT AND OF THE COUNCIL LAYING DOWN HARMONISED RULES ON ARTIFICIAL INTELLIGENCE (ARTIFICIAL INTELLIGENCE ACT), https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:52021PC0206, art. 52 "Transparency obligations for certain AI systems" includes disclosure obligations, but principally in the context of "deep fakes." The EU Parliament's amendment 487, https://www.europarl.europa.eu/doceo/document/TA-9-2023-0236_EN.pdf , proposes to broaden the transparency obligation to cover disclosure of copyrighted works: "Where the content forms part of an evidently creative, satirical, artistic or fictional cinematographic, video games visuals and analogous work or programme, transparency obligations set out in paragraph 3 are limited to disclosing of the existence of such generated or manipulated content in an appropriate clear and visible manner that does not hamper the display of the work and disclosing the applicable copyrights, where relevant."

25. If AI-generated material is found to infringe a copyrighted work, who should be directly or secondarily liable—the developer of a generative AI model, the developer of the system incorporating that model, end users of the system, or other parties?

Depending on the facts, any or all of these actors might be directly or secondarily liable. While, by virtue of the prompts she furnishes, the end-user (the person who prompts the program to produce an output) may have the most direct relationship to an infringement, upstream actors may also play a role cognizable under principles of direct or secondary liability. For example, if the model or system developer has the ability to design the model or the system to preclude certain kinds of infringing outputs (for example, not to reproduce artists' signatures), failure to design accordingly might constitute a form of "willful blindness" to the infringement.

We believe that generative AI systems do not qualify for the sec. 512 safe harbors. The systems are not host service providers "stor[ing] at the direction of a user [infringing] material that resides on a system or network controlled , , , by or for the service provider." The training data is not user-posted content, and, while end-users enter prompts, the prompts themselves are not infringing content stored at the behest of the user. The infringing content may be incorporated in

the outputs, but the system itself generates the outputs. As a result, common law principles of secondary liability would apply, and can evolve to fit the circumstances.

25.1. Do "open-source" AI models raise unique considerations with respect to infringement based on their outputs?

Assuming the question refers to open source content on which the model was trained, whether outputs incorporating portions of open source software or data infringe may depend on the terms and conditions of access to the open source content. While the Terms of Service may permit reproduction and dissemination of the code or data, there may be other restrictions, such as attribution, which, if not honored, might constitute a "material breach" terminating the license and rendering the use unauthorized and infringing.

26. If a generative AI system is trained on copyrighted works containing copyright management information, how does 17 U.S.C. 1202(b) apply to the treatment of that information in outputs of the system?

See answer to Question 24

28. Should the law require AI-generated material to be labeled or otherwise publicly identified as being generated by AI? If so, in what context should the requirement apply and how should it work?

In general, if the output contains AI-generated material for which human authorship cannot be attributed, the output as a whole, or the relevant portions, should be labelled. The Copyright Office could provide Guidelines for what and how to label, particularly when only portions of the work are AI-generated (and lacking sufficient human authorship).

Labeling serves several goals: it counters attempts to pass off "deep fakes" and spurious text, art, etc. as the "real thing," thus protecting the public against misrepresenting outputs "in the style of" as true works. It also protects performing artists against misrepresentation of their contributions to audio or audiovisual productions. It signals to the public which outputs or components of works, are not copyrightable, and therefore may be freely copied or reworked. The Copyright Office may require these disclosures in applications to register works; we are not sure the Office has the authority to require inclusion of those disclosures in copies or communications to the public of AI-generated, or AI-component productions. If that authority lies only with Congress, then Congress should also consider prohibiting some misleading AI practices altogether.

In addition, private sector endeavors to systematize AI content-labelling, such as Adobe's Content credentials,[23] should be encouraged, as should Best Practices in general.

28.1. Who should be responsible for identifying a work as AI-generated?

The person who makes the AI-generated output available to the public; that person is often the proximate cause of the output's generation, and therefore is aware of how the output was produced. If the person obtaining the output and the person who disseminates it are not the same, the labelling obligation should fall on the person who requested the output. If the disseminator knows or should have known the output was AI-generated, the obligation should attach to that person as well.

That said, there are important predicate questions. What is the legal basis of an obligation to identify a work as AI-generated? To what extent does the Copyright Office have authority to impose that obligation. Would the obligation need to be grounded in a legislative enactment?

28.3. If a notification or labeling requirement is adopted, what should be the consequences of the failure to label a particular work or the removal of a label?

Our response depends on the legal basis of the notification or labelling requirement. If the requirement is imposed through the registration process, then the remedy for non-compliance or removal would probably have to be limited to the existence or efficacy of the registration. For example, akin to sec. 401(d) ["If a notice of copyright in the form and position specified by this section appears on the published copy or copies to which a defendant in a copyright infringement suit had access, then no weight shall be given to such a defendant's interposition of a defense based on innocent infringement in mitigation of actual or statutory damages,"], but more forcefully, Copyright Office regulations might provide that if a copy to which a defendant had access did not identify the AI-generated portions, then, any or all of the following consequences will attach: (1) the registration will be cancelled; (2) the registration will be suspended pending provision of the requisite labelling information and any pending infringement action will be held in abeyance until the registration is corrected; (3) statutory damages and attorneys fees will not be available even if the work had been registered before the commission of the infringement. We believe these sanctions are Berne-compatible because art. 5(2)'s no formalities rule applies to "works . . . protected under this Convention"(art. 5(1)); under art. 2, a work must be an "intellectual creation," but if it lacks human authorship, it is not an "intellectual creation."[24]

---

[23] *See* Benj Edwards, *Adobe launches new symbol to tag AI-generated content—but will anyone use it?*, Ars technical, Oct 11, 2023, https://arstechnica.com/information-technology/2023/10/adobe-launches-new-symbol-to-tag-ai-generated-content-but-will-anyone-use-it/

[24] *See* S. Ricketson & J.C. Ginsburg, INTERNATIONAL COPYRIGHT AND NEIGHBOURING RIGHTS: THE BERNE CONVENTION AND BEYOND para 7.15 (Oxford U. Press, 3d ed. 2022).

If the source of the labelling requirement is legislative, the subject(s) of the requirement should be liable for injunctive relief (obligation to add the relevant information) as well as damages to any person harmed by the failure to label.

Additional Questions About Issues Related to Copyright

30. What legal rights, if any, currently apply to AI-generated material that features the name or likeness, including vocal likeness, of a particular person?

Currently, a variety of state right of publicity laws may apply, if they cover simulated, as well as actual visual and vocal likenesses.

Section 43(a) of the Lanham Act may apply to passing off a product or a performance as being "by" or "with" the creator or performer. Section 43(a) addresses:

> any false designation of origin, false or misleading description of fact, or false or misleading representation of fact, which—
>
> (A) is likely to cause confusion, or to cause mistake, or to deceive as to the affiliation, connection, or association of such person with another person, or as to the origin, sponsorship, or approval of his or her goods

While the Supreme Court's *Dastar* decision[25] rejects characterizing the author as the "origin" of goods, one can still argue that falsely or misleadingly attributing a product or a performance confuses or deceives as to sponsorship or approval because a creator or performer who did not approve the product or performance would not be participating in its production.

31. Should Congress establish a new federal right, similar to state law rights of publicity, that would apply to AI-generated material? If so, should it preempt state laws or set a ceiling or floor for state law protections? What should be the contours of such a right?

The recently proposed "No Fakes Act" would protect "image, voice, and visual likeness" against unauthorized "digital replicas" (and would not preempt state law).[26] It is an interesting start, but notably, includes no obligation to disclose that an individual's "image, voice, [or] visual likeness" is incorporated in a "digital replica." This NOI is probably not the place for a full analysis and critique of the "No Fakes Act", but we will be happy to provide one in reply comments, if warranted.

32. Are there or should there be protections against an AI system generating outputs that imitate the artistic style of a human creator (such as an AI system producing visual works "in the style of" a specific artist)? Who should be eligible for such protection? What form should it take?

---

[25] *Dastar Corp. v. Twentieth Century Fox Film Corp*., 539 U.S. 23 (2003).

[26] Nurture Originals, Foster Art, and Keep Entertainment Safe Act of 2023 (NO FAKES Act), S. Legis. Couns. EHF23968 GFW, 118th Cong. (2023) (draft copy).

Broadly prohibiting the generation of "in the style of" outputs goes too far; it would deprive users of an evocative teaching and learning tool, as well as of a source of amusement. But ensuring that the outputs are clearly labeled to disclose that they are imitations provides important protections against misleading and false representations. Moreover, given the permeability of the line between "style" and "expression," some "in the style of" productions may in fact infringe. Whether the claim lies in copyright or in trademarks/unfair competition, the author (human creator) of the original work should have standing to bring the action.

Some "in the style of" outputs may in fact substantially copy the source works, in a way that would not be fair use; those situations require no changes to the current copyright law. The harder scenario involves unlicensed inputting of sufficient quantities of an author's or artist's work to enable the AI generator to produce outputs that that do not reproduce expression (assuming the style/expression line can be discerned) but that threaten to replace future production of works from the creators whose works the system has ingested: Why pay an artist to create a work when one can request an image-generation program to create a similar substitute for free (or for the usage or subscription cost of the service)? But this substitution effect may not be cognizable under factor four of §107, which inquires into the effect on the market for or value of the copyrighted *work* (i.e., the works copied into the system), not the market for or value of the artist's present and future *oeuvre* in general.

By contrast, under Berne Convention article 9(2), and WIPO Copyright Treaty [WCT] art. 10, the "three-step test" authorizes member states to create exceptions and limitations to the reproduction right; the third step requires that the exception or limitation "not unreasonably prejudice the legitimate interests of the author." An AI output that competes with an artist's *oeuvre* in general or with her future work may not supplant the market for any particular copied work, but that output may indeed "unreasonably prejudice the legitimate interests of the author" in making a living and continuing her creative activities. The relevant interests in the Berne Convention and the WCT focus on the author rather than on particular works.[27] This disparity need not, however, lead to an impasse. The US fair use analysis is not limited to the four factors: courts may take other considerations into account.[28] Whether to avoid putting the US out of compliance with its international obligation to conform national copyright exceptions to the three-step test, or as a matter of purely domestic interpretation of the copyright law, US courts could look beyond §107's lack of an explicit direction to address an exception's impact on the legitimate interests of the *author*. US courts could incorporate an inquiry into that impact as an additional fair use consideration. Or Congress could amend §107 to take into account author interests, including attribution, but we are generally wary of legislative approaches in the fair use

[27] By contrast, art. 13 of the TRIPS (Trade Related aspects of Intellectual Property) Accord formulates the third step with reference to the "rightsholder," thus attenuating the authorship-based interpretation of the Berne Convention and the WCT.

[28] *See, e.g., Rogers v. Koons*, 960 F.2d 301, 309 (2d Cir. 1992) (the defendant knowingly infringed in bad faith; *see also MCA, Inc. v. Wilson*, 677 F.2d 180, 183 (2d Cir. 1981)); *Authors Guild v. Google, Inc*., 804 F.3d 202, 207 (2d Cir. 2015) (the use provides a significant benefit to the public; see also Perfect 10, Inc. v. Amazon.com, Inc., 508 F.3d 1166, 1168 (9th Cir. 2007)); *Nunez v. Caribbean Int'l News Corp*., 235 F.3d 18 (1st Cir. 2000) (the defendant attributes the plaintiff's work to the plaintiff); *Triangle Publ'ns, Inc. v. Knight-Ridder Newspapers, Inc*., 626 F.2d 1171, 1176 (5th Cir. 1980) (the use is consistent with industry practices).

context. Finally, even under a solely work-based interpretation of §107(4), one may observe that the wholesale commercial copying of an artist's works into training data in order to enable stylistically similar outputs jeopardizes not only the artist's future employment or commissions, but also devalues the actual works copied, because the image-generation program can produce outputs that compete with already-created works as well.

Standing: if the human creator is the "legal or beneficial owner" of the infringed rights, she will have standing to bring an infringement action. Standing under §43(a) of The Lanham Act is broader, reaching to "any person who believes that he or she is or is likely to be damaged by such act," but the source of authors' or artists' complaints may not lie merely in deceptive labelling.[29] Authors and artists who do not have standing to bring a copyright infringement action – authors who have granted their rights for a lump sum and therefore are not "beneficial" copyright owners, and creators of works made for hire – might nonetheless warrant protection against some forms of AI misappropriation. The Copyright Office might consider recommending that Congress extend a right of action to these creators, since "in the style of" outputs may well compete with future opportunities, including commissions to create works for hire.[30]

Respectfully Submitted,

Shyamkrishna Balganesh	Jane C. Ginsburg	Philippa S. Loengard
Faculty Co-Director	Faculty Co-Director	Executive Director

[29] 15 U.S.C. §1125(a)(1)(B).

[30] VARA does not require that a visual artist be a copyright owner, but VARA does not extend rights of attribution and integrity to artists who create works made for hire because VARA excludes works made for hire from its protections. See 17 USC sec 101 (definition of work of visual art). VARA would be very limited here, no, since it would really only apply to select photographs that were input, and even then the photo itself would not be altered (and it would be signed).