Within the timeframe, I only managed to answer a subset of all the questions that I have knowledge or opinions for. If allowed, I will add my answers to other questions during the replies time window.

2. Does the increasing use or distribution of AI-generated material raise any unique issues for your sector or industry as compared to other copyright stakeholders?

Yes. I work as a senior concept artist in the gaming industry during the day and as an independent artist/illustrator at night, so my answer will be divided into different sections:

I. **Within game studios:**

Companies including the one that I work in are starting to train AI models based on our (work-for-hire) artwork with the intention of specifically mimicking our art styles. This is done for the potential of quicker and/or mass production of artwork that fits into our project aesthetics, but it is deeply concerning to me in many ways.

1. **Mimicry of artistic identity**

Stylistic mimicry can lead to problems for artists on a personal level just like deepfake does to people in real life. When a human artist produces an artwork, even when we try to adjust our styles to fit the needs of a project, traces of our personal habits and stylistic signatures will remain and the people familiar with them can often tell who painted which image. In this sense, our art styles are linked to our identities in people's perceptions, and having an AI model of that, which the company has full control over, is both risky and disturbing to us as individuals. We signed up for our jobs before the prevalence of generative AI to produce individual art pieces, not to sell our artistic identities to the companies. However, we currently have no legal grounds to object to this.

2. **Job displacement at a cost to the public**

These AI models have the potential to permanently displace a large portion of artists from our job positions. While some may argue that job displacement as a result of technological innovation is a natural process, the problem this time is that it will also downgrade the quality of entertainment the public receives. Even though AI generations cannot reach the bespoke quality of human creations, the ease and low cost of using AI still makes it a desirable option for businesses from a profitability standpoint. More importantly, in the long term, I foresee a decrease of new artists entering this field due to bad career prospects. This is problematic because while generative AI is good at producing "good enough" quality images at a large scale, it does not have the ingredients to push boundaries, as it is fitted on existing content. (Just to clarify, if done well, it is possible for generative AI to produce more intricate details and finer

renderings, or even to come up with new aesthetics through combining existing ones, but it does not go beyond the status quo in the “interpretive” aspect of artistic expressions. For instance, an AI model trained on realism will not be able to generate impressionistic artwork, as this leap of visual interpretation requires the mind of a human.) The entertainment space will likely be stagnantly filled with profitable but generic content unless we encourage new human artists with their own voices and interpretations to come into the scene.

II. Among indie artists and in the general creative ecosystem:

Many studio artists like myself also practice art after hours, either posting personal work online or selling them in marketplaces. This is the opportunity for us to project our creative voices to the world, market ourselves and inspire each other. It is also through this avenue that I get to experience the lives of indie artists and freelancers, who are also facing problems caused by generative AI.

1. Parasitic behavior and performance punishment:

Within the last year, I have had multiple occasions where my artwork was mistaken to be AI generated both online and in person, possibly due to the aesthetics resembling certain classes of AI images at first glance. This is disheartening to me as I put in tremendous effort telling a human story through each image. One may suggest I look for alternative art styles, but the problem is, it takes years for an artist to develop a new style, but fine tuning an AI model to fit that style can take as little as 10 images. In this sense, the technology works like a parasite. The further irony is that the more appealing an artist’s body of work is, the more likely these works will be used as training materials. Artists like myself do not know what to look forward to because putting in more effort honing our craft only means a higher chance of getting our personal voice appropriated. This kind of performance punishment is detrimental to the creative community.

2. Devaluation of existing art styles through overexposure:

Before generative AI, this was a negligible issue that may have even played a positive role in transforming the creative landscape over time, but it is now massively amplified by generative AI to a point that it is disruptive and unsustainable. In general, for an artwork to appeal to the public, it needs to have some level of familiarity to the audience to form a connection, and an element of novelty that piques interest. Without this balance, the work becomes either too generic or too difficult to grasp, thus limiting its commercial success. This applies to both subject matters and aesthetics. The problem with generative AI is that it latches onto a successful art style and massively produces similar images that flood the content space, quickly overwhelming the audience with this

aesthetic and devaluing it. Based on my experience, many people turn away from an artist's work if they recognize a similar style in AI images. This significantly harms the artists whose works are heavily influential in AI training, even if the AI generations are neither directly plagiarizing (by current standards) nor used for commercial gains by the prompters.

3. **New artists drowned and established artists start to paywall:**

Due to its high output, the amount of AI generation in the content space will likely exceed human creations in the future. Finding the seed to exposure for a new artist was already a difficult task before, and it will only get worse onwards. On the other side of the spectrum, artists with a large enough follower base are incentivized to paywall artwork that would otherwise be free for public viewing, out of fear of having them scraped into training databases. Taken together, not only do these factors obstruct the flow of quality information that the creative community relies upon for inspiration, but also weaken the human connection that arts and culture are meant to bring to the general public.

III. As a consumer:

Most artists are also consumers of creative content. Personally, I also interact directly with consumers in marketplaces, and because of this, I can speak from their perspective on issues related to generative AI.

1. **"Bad money drives out good"**

As mentioned previously, my artwork tends to be quickly dismissed by those who mistaken them to be AI generated. This is because many consumers perceive AI generated content to hold less value than human creations, due to their lack of (perceived) association with skill, effort and human connections. However, generative AI is now capable of producing images that are difficult for the layperson to tell without careful examination. One recent example is when the voice actress Grey DeLisle brought an art print of her voiced character, Daphne, only to find out on social media that it was AI generated and demanded refunds. This is a problem that is increasingly plaguing art marketplaces. Without a way for people to pick out authentic products from lookalikes, consumers may lose trust in the current forms of creative content overall.

2. **Loss of authorship information and origin stories**

When an artwork piques my interest, I would like to know the author and story behind it, what the sources of inspirations are and how artistic choices are made. Not only does this pay respect to the author, looking into the experience of another person is a fundamental part of human learning and helps us form empathic relationships with the

world. However, AI generations do not contain this information. Due to their statistical fitting nature, rather than showing how individual humans express themselves, AI outputs are more akin to a general statement of “this is how others express themselves”, yet no information is given toward who the “others” are nor why they make such expressions.

I would like to illustrate this point with a simplified scenario. Say, for example, if two artists were asked to paint “a neighborhood on a summer day”, one of them might relate to sensory experiences such as heat, bright sunlight, and the scent of summer vegetation, hence painting a warm, saturated and overexposed street scene. Another artist might relate back to her childhood during summer vacations and paint a cheerful imaginary neighborhood playground. Both of these are individual expressions that embody each artist’s personality, habits, life story and world interpretation. Now, if a program takes the two images as data points and creates an “average” of the two, the resultant image, despite still looking like an artist’s creation following the same prompt, obscures the original intentions yet provides very little added value. Of course, this is an oversimplification of how generative AI works, but the objective is the same: to create a fitted model approximating existing data, and generate new data based on interpolation. In the simplified scenario, one can still trace back to the original artists, but when the dataset extends to all the content available on the internet, tracing becomes impossible, yet for any set of existing human data, there can be an infinite number of “generative” interpolations that confuse the audience. As a consumer, I do not look forward to exploring a content space like this.

3. Please identify any papers or studies that you believe are relevant to this Notice. These may address, for example, the economic effects of generative AI on the creative industries or how different licensing regimes do or could operate to remunerate copyright owners and/or creators for the use of their works in training AI models. The Office requests that commenters provide a hyperlink to the identified papers.

AI Art and Its Impact on Artists
https://dl.acm.org/doi/10.1145/3600211.3604681

5. Is new legislation warranted to address copyright or related issues with generative AI? If so, what should it entail? Specific proposals and legislative text are not necessary, but the Office welcomes any proposals or text for review.

Yes. For visual artists, their art style is an embodiment of their personality, taste, and habits of expression that are individualistic and signatory. It is often perceived to be the artist’s identity in

the creative space. However, unlike an actor's physical image or a singer's voice, a visual artist's style is protected neither as part of the artist's personality rights nor under copyright law, and this often leads to abuse by users of generative AI.

Before the rise of AI, stylistic mimicry was not a major issue because human artists have very little incentive to mimic the styles of others, as the time and effort needed for successful mimicry can better be spent finding and honing the artist's own creative voice. However, generative AI makes it possible for anyone to use the style of another artist with ease. This is problematic for the original artist both economically—the creative voice that took them years to develop is being devalued into a "stock"—and personally—their identity in the artistic space can now be associated with projects that they didn't approve of.

Some are suggesting making art styles copyrightable. This would cause many problems practically. For example, in collaborative projects such as game development, multiple artists have to match styles to a certain extent for consistency. (Note that, as opposed to the mimicry scenario I mentioned above, this kind of style matching is done with permission and is rarely precise.) Another problem with copyrighting art styles is that large corporations can find ways to achieve style monopolies through hiring human artists and asking them to assign rights to the company. This is not good for the creative economy overall and will destroy the livelihoods of artist employees when they leave the company.

I believe art styles should be protected in a separate category that specifically addresses the problems introduced by generative AI. I do not work in law, so I can only speak of a human's natural instincts. Intuitively speaking, there are different "degrees" of personal identifiers. Conventional ones, such as a person's name and physical likeness, are directly linked to the individual's identity and should be protected by publicity rights. At the same time, there is another category of identifiers that, while not rigorously tied to one's physical identity, "suggests" authorship in public perception and can still harm the individual when mimicked and misused—a problem that is now exacerbated by machine learning. This includes a person's handwriting, an artist's style, and other personal habits of creative expression. This category of characteristics is a mixture of personal identity and intellectual accomplishment, and should be protected in a special way. Because art styles can be ambiguous (two artists can naturally draw similarly to each other) and changing over time (an artist's style may evolve), one cannot simply assign a certain "look" to an artist, but rather, a better way of protecting art styles is to forbid certain actions from being performed on them. Specifically, these are actions introduced in the age of AI that are style-extracting and style-devaluing in nature, including but not limited to 1) directly using an artist's work for machine learning without permission (even if it is mixed with other images), 2) having a human mimicking the style of another artist for the purpose of machine learning, and 3) using an AI model trained on an artist's style for a purpose or duration of usage different than what was agreed upon. None of these affect human artists learning from and influencing each other, as I believe any new legislation should not have to majorly alter the way creatives behave when AI is left out of the equation.

Regarding getting permission from copyright holders for AI training, I believe it is equally, if not more, important to get consent from the “physical” authors of the artwork who may have signed away their rights to that particular artwork. In the broader sense, style extraction can do even more harm to an artist than plagiarism or distortion of individual copyrighted pieces, and should be addressed seriously in order to have a sustainable creative ecosystem. I believe there should be some unwaivable form of truth connecting an artwork to the physical author whose human imprint is contained within the work, so that creatives can object to having their styles extracted through machine learning regardless of whether or not they are the copyright holders (or whether or not they waived the moral rights to their creations).

Another possible scenario is that in an employment relationship, employers will try to have employees sign away all their rights and agree to have their work used for AI training. I believe blanket terms asking for generic and perpetual usage of an artist’s style in an employment contract should be outlawed. This is due to its potential to cause irreversible damage to the artist’s livelihood, even after the employment relationship ends, through a single step of negligence or in a scenario of power imbalance in a bad economy. Rather, permission should only be obtained with clear statements of the purpose and the duration of usage of the AI model to be trained, and that the model should be destroyed once the term ends.

9. Should copyright owners have to affirmatively consent (opt in) to the use of their works for training materials, or should they be provided with the means to object (opt out)?

Opt-in for three reasons:

1. The sheer number of AI models being trained at any given time makes it impossible for copyright owners to track where their work is being used. One can confirm this by examining the website civitai.com.
2. Most of the time, copyright owners do not know that their work is being trained on until the model is released. By then it is already too late because AI models cannot "unlearn".
3. Some copyright owners may not be alive.

9.3. What legal, technical, or practical obstacles are there to establishing or using such a process? Given the volume of works used in training, is it feasible to get consent in advance from copyright owners?

For fine-tuning, which is usually targeted at specific sources and only requires from ten to a few hundred images, it is definitely feasible.

For pre-training, this may be difficult, but I believe it should be up to the company to figure out a way to do so if they wish to engage in such businesses to start with. It is not a good argument to say “I’m going to harm someone for my own gains because it’s too difficult otherwise.”

9.5. In cases where the human creator does not own the copyright—for example, because they have assigned it or because the work was made for hire—should they have a right to object to an AI model being trained on their work? If so, how would such a system work?

The creator should have a right to object. Please see my answer to question #5.

15. In order to allow copyright owners to determine whether their works have been used, should developers of AI models be required to collect, retain, and disclose records regarding the materials used to train their models? Should creators of training datasets have a similar obligation?

Yes. Both developers of AI models and creators of training datasets should be required to collect, retain and disclose records regarding training materials.

15.1. What level of specificity should be required?

The level of specificity should be enough for individual copyright holders to identify whether or not their work is included in the training dataset.

15.2. To whom should disclosures be made?

Disclosure should be made to the public because large datasets can potentially contain anyone’s work. It should also be presented in a navigable way for consumers to view, as it holds a similar purpose as a product label or movie credits.

15.4. What would be the cost or other impact of such a recordkeeping system for developers of AI models or systems, creators, consumers, or other relevant parties?

Individual creators do not have the resources to always be kept up to date about which dataset contains their work, and there are very limited actions they can take if they do spot their work in an existing dataset. Therefore, a recordkeeping system alone is insufficient, and will have to

work in conjunction with another system that obtains permission from the copyright holders before training takes place.

16. What obligations, if any, should there be to notify copyright owners that their works have been used to train an AI model?

Developers of AI models should be obliged to seek permission from copyright owners before using their work as training material (thereby notifying copyright owners).

28. Should the law require AI-generated material to be labeled or otherwise publicly identified as being generated by AI? If so, in what context should the requirement apply and how should it work?

Ideally, yes. Access to truthful information is the foundation on which a functional democratic society operates. It is also important to foster consumer trust in the creative industry. However, I don't know if such an effort is technically feasible.

28.1. Who should be responsible for identifying a work as AI-generated?

A truthful and thorough labeling system might only be achievable through a combined effort of the provider of the AI service, users that generate and/or distribute AI output, an auditing system and community tagging. I believe the AI service provider and the user should bear most of the responsibilities.

28.2. Are there technical or practical barriers to labeling or identification requirements?

The boundary between human creation and AI-assisted content production is getting increasingly blurry, especially as popular software such as Adobe Photoshop incorporates generative features. If an image only contains 10% AI-generated content because the user chose to extend the canvas of a photo for a better composition, would this image be considered an AI generation? What about an artist that started with an AI-generated image and painted on top of it, eventually covering the entirety of the canvas so that the only AI-relevant aspect is the starting composition? What about a character design sheet for a game or animation that is entirely drawn by a human, but the character *design* itself is taken from an AI generation? Any labeling or identification will need to specifically identify which parts of the image are AI-generated, or which stage of the production process was AI involved, rather than using simple "AI" or "no AI" tags.