I'm a digital illustrator, storyboard artist, and undergraduate student of Political Science. I have been drawing my entire life, but started seriously practicing art about 8 years ago. As such, most of my answers will refer to text-to-image AI generators unless otherwise specified.

2) Does the increasing use or distribution of AI-generated material raise any unique issues for your sector or industry as compared to other copyright stakeholders?

The increasing use and distribution of AI generated materials has and will continue to negatively impact visual artists without intervention from the Copyright Office to protect original creators. AI image generators such as Stable Diffusion rely on the ingestion of an artist's entire lifetime of works, with some products specifically advertising artists' names as a prompt for users, without consulting the artist first. As Jiang et al. (2023)'s paper "AI Art and its Impact on Artists" details, the visual arts face unique consequences such as economic loss, forgery, reinforcement of harmful stereotypes, and chilling effects for cultural production (https://dl.acm.org/doi/10.1145/3600211.3604681).

To expand on economic loss: one high profile case Jiang et. al (2023) references is Marvel Studios' use of AI generated imagery for the opening credits of their TV series *Secret Invasion* . In contrast to the studio's past TV series in which numerous artists and illustrators are credited for the making of the title sequence, Secret Invasion only lists one (Jiang et al. 2023). Other examples include large game studios such as Niantic and Rayark, which have been accused of using AI generated assets after mass layoffs of studio artists (https://kotaku.com/pokemon-go-adventures-abound-art-ai-artificial-niantic-1850793819)(https://kotaku.com/ai-art-layoff-video-game-studio-pc-midjourney-aigc-1850489333).

9) Should copyright owners have to affirmatively consent or "opt in" to the use of their works in training materials, or is it enough to be provided the means to "opt out" of training?

Copyright owners must have to affirmatively "opt in" to the use of their works for training materials, because "opt out" systems are too burdensome and complex for individuals to meaningfully make use of. In the first place, if users aren't aware of any such policy, opt-out de facto punishes users that may otherwise want their copyrighted material removed. If they are aware of such a policy, the process can impose large time costs - for an artist to opt out of Stable Diffusion, they need to know which of their pieces are in the model, then upload each individual art piece, go through every search result, and select matches to their piece. This process is completely unrealistic for individual artists to pursue, especially with how common it is to have catalogs of hundreds and thousands of drawings. In the case of Meta's opt out request tool, it is

functionally impossible for one to delete their data without providing personal information, and even then the company can refuse to process the request for any reason (https://www.wired.com/story/meta-artificial-intelligence-data-deletion/). As Kapoor and Narayan point out, the most significant utilizations of Stable Diffusion's opt-out policy came from organized collectives: "Of the 80 million images opted out, only about 40,000 came from individual artists. That's less than 1 in 2,000. The rest were opted out in bulk by platforms such as Shutterstock and Artstation." (https://www.aisnakeoil.com/p/artists-can-now-opt-out-of-generative) Rights should not be accessible only to those with the time and resources, and regulations should reflect this reality through an "opt-in" requirement. As The Atlantic notes: "Current AI models have already digested a massive amount of work, and even if a piece of art is kept away from future programs, it's possible that current models will pass the data they've extracted from those images on to their successors"(https://www.theatlantic.com/technology/archive/2023/10/openai-dall-e-3-artists-work/675519/). It is critical that "opt-in" is the standard going forward to prevent the spread of further violations and protect the rights of original creators.

15) In order to allow copyright owners to determine whether their works have been used, should developers of AI models and/or creators of training datasets be required to collect, retain, and disclose records regarding the materials used to train their models?

It is essential that developers of AI models and training sets are mandated to collect, retain, and disclose records on what materials were used to train their models. Developers are best poised to organize mechanisms for record collection, retention, and disclosure, although a third party would be needed to ensure compliance. Anything less would make it near impossible for individual copyright holders to determine when their works have been used.

18) Under copyright law, are there circumstances when a human using a generative AI system should be considered the "author" of material produced by the system? Is selecting what material an AI model is trained on and/or providing an iterative series of prompts sufficient to claim authorship of the resulting output?

Under the current AI models that leave human users with little to no control over the end product, it should not be possible for a human to be considered an "author" of AI generated material.

It is true that some models allow you to train them using a selection of specific works and prompts. It is also true that the specific selection and coordination of individual works to create a collective work, like a collage, can be protected under US Copyright Law (https://copyright.gov/eco/help-collective-work.html). However, this does not mean that a human can gain authorship of an AI generated image merely for selecting training materials, as the final image is not the selection or arrangement of the

materials themselves, but rather a completely separate image. Therefore, mere selection of training material is not sufficient enough to claim authorship.

Even if in combination with iterative prompting and other methods like weighing to retain and replace specific parts of an AI generated image, it does not change the fact that most of the output is due to the AI's predictive algorithm. Therefore, the specific expression of an idea/prompt is due to the AI, and not the human user, which disqualifies them from a claim of authorship.

21) In the Constitution, the justification behind the Copyright Clause is to "promote the progress of science and useful arts." Does this clause permit copyright protection for AI-generated material, and would such protection promote the useful arts?

The Copyright Clause does not permit the protection of solely AI-generated material. The ingestion, retention, and appropriation of artist' works to create biased and inferior-but-good-enough AI products to ultimately replace them does not promote the useful arts, but rather degrades it. I know artists that have stopped sharing works, even taking down archives of their work, to prevent their artistic creations from being appropriated by AI. I have heard from students and professionals in the arts and entertainment feeling discouraged from pursuing it as a career path, needing to take second jobs to make up for lost income, and considering leaving the profession altogether as companies see AI as a good enough cost cutting measure. Decreasing the viability of art as a career path for humans makes our visual culture less rich and less diverse, especially when AI is known to amplify biases, reinforcing a rigid view of how our world should be (https://www.businessinsider.com/ai-image-prompt-for-african-workers-depicts-harmful-stereotypes-2023-4). For example, Bloomberg found that Stable Diffusion exhibited strong bias for white people and for men when prompting the AI with various occupations. Higher paying jobs were dominated by subjects with lighter skin tones and while men outnumbered women in high paying occupations (https://www.bloomberg.com/graphics/2023-generative-ai-bias/). A study by Luccioni et. al (2023) corroborated these biases towards whiteness and masculinity across DALL-E2 in addition to Stable Diffusion (https://arxiv.org/abs/2303.11408).

28) Should the law require AI-generated material to be labeled or otherwise publicly identified as being generated by AI? If so, in what context should the requirement apply and how should it work?

The law should mandate AI-generated images to be disclosed as such through clearly identifiable means at the point of generation, like a watermark, or a disclosure provided by digital platforms hosting these images, such as Google's "AI-Generated" label. These methods are necessary to combat scams in which customers such as Grey DeLisle commission a piece in the range of hundreds of dollars with the expectation it

was a human "original", and are instead provided with AI generated art (https://thechainsaw.com/nft/nft-art/scooby-doo-daphne-voice-actor-ai-art-scam/). In other cases, such scams can be compounded with theft of an original artist's sketch to produce the AI generated output (https://thechainsaw.com/nft/scammers-artist-draft-ai-art-resell-japan/). Outside of monetary harms, disclosures are necessary to combat the appropriation of an artist's style to promote harm. For example, famous comic artist Sarah Andersen details how her works have been re-edited by Neo-Nazis to promote Holocaust denialism and then spread across the internet (https://www.nytimes.com/2022/12/31/opinion/sarah-andersen-how-algorithim-took-my-work.html). Without disclosure, it is easy to imagine how an AI can amplify such reputational harms to a creator, potentially costing them jobs, income, careers.

Outside of the arts, disclosures on all generative AI media are key to fighting political misinformation perpetuated by fake AI-generated photos. The outcome of Slovakia's 2023 elections was potentially swayed through a deepfaked audio recording of a prominent party leader discussing how to rig elections, allowing an opposing party that campaigned on withdrawing military aid to Ukraine to win a tight election (https://www.wired.com/story/slovakias-election-deepfakes-show-ai-is-a-danger-to-democracy/). The increasing realism of image generators will soon allow any state, party, or individual to generate inflammatory photos to stoke conflict or otherwise mislead the public, especially during times of crisis. This can increase animosity towards a specific ethnic group, providing governments with the popular support they need to enact and cover up their atrocities, much like how Myanmar brewed support for the Rohingya genocide through massive misinformation campaigns on Facebook (https://www.amnesty.org/en/latest/news/2022/09/myanmar-facebooks-systems-promoted-violence-against-rohingya-meta-owes-reparations-new-report/). Though companies like Google are already working on tools to allow for disclosure, a law could outline standardization and promote widespread adoption of AI disclosures, which will help arm users with information on whether an image and associated narrative is true or not. These disclosures should comport with free speech principles, as combating false or hateful speech by providing more speech in the form of disclosure is content neutral and narrowly tailored.