Comment to U.S. Copyright Office - Artificial Intelligence and Copyright

Ed Newton-Rex

I recently resigned from my position as VP of Audio at Stability AI, because I disagree with the position they take on fair use outlined in their response to the US Copyright Office, and with the similar position held by many other AI companies. I've been working in generative AI since 2010, when I founded one of the first generative AI startups. Here, I include a statement I made when I resigned from Stability AI, as well as some further thoughts on copyright and generative AI.

Statement on my resignation

Originally published in Music Business Worldwide.

I've resigned from my role leading the Audio team at Stability AI, because I don't agree with the company's opinion that training generative AI models on copyrighted works is 'fair use'.

First off, I want to say that there are lots of people at Stability who are deeply thoughtful about these issues. I'm proud that we were able to launch a state-of-the-art AI music generation product trained on licensed training data, sharing the revenue from the model with rights-holders. I'm grateful to my many colleagues who worked on this with me and who supported our team, and particularly to Emad for giving us the opportunity to build and ship it. I'm thankful for my time at Stability, and in many ways I think they take a more nuanced view on this topic than some of their competitors.

But, despite this, I wasn't able to change the prevailing opinion on fair use at the company.

This was made clear when the US Copyright Office recently invited public comments on generative AI and copyright, and Stability was one of many AI companies to respond. Stability's 23-page submission included this on its opening page:

"We believe that AI development is an acceptable, transformative, and socially-beneficial use of existing content that is protected by fair use".

For those unfamiliar with 'fair use', this claims that training an AI model on copyrighted works doesn't infringe the copyright in those works, so it can be done without permission, and without payment. This is a position that is fairly standard across many of the large generative AI companies, and other big tech companies building these models—it's far from a view that is unique to Stability. But it's a position I disagree with.

I disagree because one of the factors affecting whether the act of copying is fair use, according to Congress, is "the effect of the use upon the potential market for or value of the copyrighted work". Today's generative AI models can clearly be used to create works that compete with the copyrighted works they are trained on. So I don't see how using copyrighted works to train generative AI models of this nature can be considered fair use.

But setting aside the fair use argument for a moment—since 'fair use' wasn't designed with generative AI in mind—training generative AI models in this way is, to me, wrong. Companies worth billions of dollars are, without permission, training generative AI models on creators' works, which are then being used to create new content that in many cases can compete with the original works. I don't see how this can be acceptable in a society that has set up the economics of the creative arts such that creators rely on copyright.

To be clear, I'm a supporter of generative AI. It will have many benefits—that's why I've worked on it for 13 years. But I can only support generative AI that doesn't exploit creators by training models—which may replace them—on their work without permission.

I'm sure I'm not the only person inside these generative AI companies who doesn't think the claim of 'fair use' is fair to creators. I hope others will speak up, either internally or in public, so that companies realise that exploiting creators can't be the long-term solution in generative AI.

Further to this, I wanted to address a few other points I've seen in companies' responses to the copyright office, and in public discourse.

1. Isn't AI training just the same as humans learning? Why should copyright only apply to AI learning?

AI training is entirely different to human learning because AI's output easily scales. A single AI system could theoretically produce enough creative output (e.g. background music) to significantly reduce the demand for all prior output in that market. This has never been possible with human learning. Moreover, human creators have, for a long time, entered the creative market with a social contract in place that expects future competitors to learn from their output. This is understood by all who enter the market. On the contrary, a highly scalable, automated system training on their output has *not* been part of the social contract they have knowingly entered.

2. Does licensing training data shut out smaller AI developers?

No. While some licensing models involve large upfront payments, others do not. One simple model that can work for any size developer is a revenue share between the content rights-holder and the AI provider, which can be achieved without any upfront payment. Small teams and small companies are already putting in place such models, disproving the argument that they will be shut out by licensing.

3. Is licensing training data infeasible because of the amount of data you need?

Not necessarily. In some modalities (e.g. image and music), various teams have demonstrated that you can get very good results from datasets that are far smaller than are traditionally thought necessary for generative AI models. These models have been used successfully in commercial offerings. It is true that, in some modalities (e.g. text), you still

need a very large amount of data to train the best models. But it is by no means certain that this will always be the case.