

**Before the
UNITED STATES COPYRIGHT OFFICE
Washington, D.C.**

In the Matter of:

Artificial Intelligence and Copyright

Docket No. 2023–6

**REPLY COMMENTS OF
AMERICAN ASSOCIATION OF INDEPENDENT MUSIC,
RECORDING ACADEMY, AND
RECORDING INDUSTRY ASSOCIATION OF AMERICA, INC.**

The American Association of Independent Music (“A2IM”), the Recording Academy and the Recording Industry Association of America, Inc. (“RIAA”) are pleased to provide these reply comments in response to the Notice of Inquiry (“NOI”) published by the Copyright Office (the “Office”) on August 30, 2023. *See* 88 Fed. Reg. 59942.

A2IM and RIAA are trade associations that represent hundreds of recorded music companies of all sizes. Our members create, manufacture, and/or distribute sound recordings representing a large majority of legitimate recorded music consumption in the United States.¹ The Recording Academy (“Academy”), best known for celebrating artistic excellence through the annual GRAMMY Awards, is a leading trade association for individual music professionals, representing nearly 20,000 artists, performers, songwriters, composers, and studio professionals across all genres and regions.

Introduction and Executive Summary

As described in more detail in the initial comments to the NOI (“Initial Comments”) filed by each of our groups as well as others,² the music industry has been using various forms of artificial intelligence (“AI”) for years to enhance and support human creativity. As an industry, we are currently working hard to harness the enormous promise of AI to support our artists and songwriters, while simultaneously pushing back against the dangers posed by AI through unlawful exploitation of our members’ works and their artists’ voices and likenesses.

New uses of AI in the music space are unfolding at a rapid pace. Indeed, since the filing of Initial Comments, several new music-related AI ventures have been announced. On November 16, for example, YouTube announced an experimental feature that is enabling creators to use AI

¹ For a fuller description of A2IM and RIAA, *see* Initial Comments of A2IM and RIAA (January 5, 2023) (“A2IM-RIAA Initial Comments”).

² *See, e.g.*, Initial Comments of Public Knowledge at 17 (“Music in particular has had a long relationship with GAI: voice-synthesizer Vocaloid, released in 2007, is so popular that the app’s AI-powered cartoon avatar has gone on world tours annually since 2014; David Bowie used a custom-built “Verbasizer” app to generate song lyrics as early as 1995; and Adaptiverb, which uses machine learning to craft the perfect bespoke reverb tail has been commercially available since 2016”) (citations omitted).

to generate new music featuring vocal simulations of well-known artists for use in clips on YouTube Shorts. Participating artists include Charlie Puth, John Legend, Sia, T-Pain, Demi Lovato, Troye Sivan, Charli XCX, Alec Benjamin, and Papoose, each of whom has approved their inclusion in this experiment.³ On November 17, CreateSafe, the music tech studio that collaborated with Grimes on an AI using a simulation of her voice, announced that it had raised \$4.6 million to launch TRINITI, a new “artistic intelligence” platform.⁴ And on November 21, the Orb and David Gilmour launched the “Metallic Spheres in Colour” AI Global Remix Project, in collaboration with Sony Music Entertainment, Legacy Records, and Vermillio, to let fans create AI remixes and personalized cover art of portions of that album.⁵

While we recognize that the nature of generative AI sets it apart from prior technological advances in a number of respects, it does not obviate the need for parties to obtain licenses for their use of copyrighted works. Indeed, the need for new businesses made possible by new technology to obtain licenses to large catalogs of creative works, in a manner that respects and supports human creators, is not new.⁶ Record companies have successfully supported artists by licensing their full catalogs to an array of “new technology” digital services – including streaming, user-generated content (“UGC”), social media, fitness apps, and more – for decades, and by licensing individual sound recordings for sync uses and for use as samples in other sound recordings. Clearly, record labels possess the know-how, the experience, and the infrastructure necessary to make the voluntary licensing of AI tools both possible and practicable.⁷ Suggestions by AI companies and others to the contrary are simply false.⁸

³ See, Elias Light, *YouTube Creators Can Now Use AI to Imitate Pop Stars*, Billboard (Nov. 16, 2023) (“YouTube’s feature, called Dream Track, creates pieces of music — voice along with musical accompaniment — based on text prompts that are up to 30 seconds in length. For now, around 100 U.S.-based creators will have Dream Track access.”) https://www.billboard.com/pro/youtube-ai-tool-sound-like-demi-lovato-troye-sivan/#recipient_hashed=7c6a8617fd97d8d9b3bedb2dfe47f5beaa94361c1c0b8f49b261a4112e0e0769&recipient_sal_t=378d9fc5d6ab6bac91b045fb06d1d5457b71aef4c528f654ecc7eaea96d2b5f.

⁴ Ashley King, *Music Tech Studio Behind Grimes AI Announces \$4.6 Million Fund for New Platform*, Digital Music News (Nov. 16, 2023) https://www.digitalmusicnews.com/2023/11/16/createsafe-music-funding-ai-platform/?utm_source=Daily+Snapshot&utm_campaign=f263ee6b19-Daily_Snapshot_November_17th_2023&utm_medium=email&utm_term=0_-77cac1fdfd-%5B%5D_EMAIL_ID%5D&mc_cid=f263ee6b19&mc_eid=2a1e77a7bc (“TRINITI is a comprehensive ‘GenAI’ music platform that aims to power creation, publishing, administration, distribution, and marketing of music, providing artists and their teams unprecedented power to generate ideas and bring them to the public.”).

⁵ *The Orb and David Gilmour Launch ‘Metallic Spheres in Colour’ Global Remix Project for Fans*, Sony Music Entertainment (Nov. 21, 2023) <https://www.sonymusic.com/sonymusic/the-orb-and-david-gilmour-launch-metallic-spheres-in-colour-ai-global-remix-project-for-fans/>.

⁶ In fact, a recent global consumer study found that fans agree that human creators should be respected. The study found that 76% of the respondents feel that an artist’s music and vocals should not be used or ingested without permission, and 74% of the respondents agreed that AI should not be used to clone or impersonate artists without permission. See *Music fans worldwide believe human creativity essential in time of AI*, IFPI, Nov. 27, 2023, <https://www.ifpi.org/music-fans-worldwide-believe-human-creativity-essential-in-time-of-ai/>.

⁷ Critically, licensing revenues are not ancillary revenues, but rather the main source of revenue for record labels and the key source of royalties for artists.

⁸ As discussed in the A2IM-RIAA Initial Comments at 7-8, potential licensees can — and do — obtain access to virtually the entire catalog of recorded music by entering into licensing agreements with the three major record companies as well as individual independent labels and aggregators/distributors.

As the Office sifts through the roughly 10,000 initial comments it received in connection with this study, we urge it to make clear in its forthcoming Report(s)⁹ that AI developers – like other innovative industries before them – can and must respect the constitutionally-based copyright interests of human creators. The Office should reject the false and dangerous dichotomy posed by some of the tech companies and their supporters in their initial comments (together, the “AI Company Comments” or “AI Company Commenters”) suggesting that humanity and, by extension, the Office and Congress, is faced with a binary choice to either reap all the benefits that AI promises *or* respect the long-established legal rights of human creators.¹⁰ There is no reason that implementations of AI cannot *both* reach their full potential *and* protect and respect the rights of creators at the same time. In fact, the application of a clear and predictable legal framework to AI is likely to foster, rather than hinder, its development by eliminating legal and business uncertainty and limiting the risk of litigation, and, most important, ensuring that our society continues investing in human creativity.¹¹

To this end, we urge the Office, in its forthcoming Report(s), to affirm that existing copyright laws and doctrines (so far)¹² are adequate to address the questions being raised by the advent of generative AI and to recommend a handful of non-copyright legislative proposals that would augment the protections offered by the Copyright Act. More specifically, we encourage the Office to do the following:

- Reject – in the strongest terms possible – the contention that the unauthorized ingestion of copyrighted works by AI systems is always (or almost always) fair use. Not only is the notion of a “*per se* fair use” rule inimical to longstanding copyright jurisprudence, as explained in detail in Section I below, the arguments advanced in support of this contention are not borne out by the cases they cite. The Report(s) should make clear that fair use is a fact specific inquiry and that, in particular, where an AI system is designed to produce outputs that compete with the unlawfully ingested works, there is a high likelihood that the use does *not* qualify as fair use.¹³

⁹ We assume from the recent Executive Order that the Office will be issuing a written report(s) in response to this NOI. Exec. Order No.14110, Section 5.2(c)(iii), 88 Fed. Reg. 75191, 75207 (Nov. 1, 2023) (hereinafter “Executive Order”).

¹⁰ See, e.g., Chamber of Progress Initial Comments at 2-4; Andreessen Horowitz Initial Comments at 2-4; Consumer Technology Association (“CTA”) Initial Comments at 1; Engine Initial Comments at 2-3; Microsoft Initial Comments at 4-5, 9; OpenAI Initial Comments at 1-5, 13-14; TechNet Initial Comments at 5-6.

¹¹ Without strong respect for the rights of creators/copyright owners, the pipeline of quality works on which AI models depend could dry up. See A2IM-RIAA Initial Comments at 20; Music Community Comments in response to the July 7, 2023 Office of Science and Technology Policy Request for Information on National Priorities for Artificial Intelligence at 11 (section on additional input), <https://www.regulations.gov/comment/OSTP-TECH-2023-0007-0231>.

¹² See A2IM-RIAA Initial Comments at 11 (noting the possibility that legislative reform may be necessary depending on how ingestion technology develops in the future).

¹³ It is disingenuous for the AI Company Commenters to claim that investors’ expectations regarding the applicability of the fair use defense to the development of AI tools have been a “critical factor in the enormous investment of private capital into U.S.-based AI companies.” Andreessen Horowitz Initial Comments at 6. To the contrary, investors have poured “billions and billions of dollars,” *id.*, into the development of AI technologies without a single court ruling on AI companies’ fair use arguments. One need only follow the money to see that investors are highly motivated to invest in this technology, notwithstanding the very real risk of a pro-copyright ruling on the fair use issue.

- Express support for federal legislation that creates a uniform, nationwide baseline to ensure that AI systems cannot unfairly exploit artists' expression, creative contributions, voices, and likenesses without their consent.¹⁴
- Reaffirm the position set forth in the *Thaler*,¹⁵ *Kashtanova*,¹⁶ and *Allen*¹⁷ cases that only works of human authorship are entitled to copyright protection, not synthetic content output by an AI system based on prompts.
- Express support for non-copyright legislation that creates mandatory recordkeeping and disclosure standards so that copyright owners will have access to the data they need to enforce their rights.
- Express support for legislation to create an administrative subpoena process (and evidentiary presumption) as outlined in the A2IM-RIAA Initial Comments.
- Focus on market-based licensing solutions and reject government-driven approaches such as statutory or extended collective licensing for AI developers. As the Initial Comments make clear, neither the would-be licensors nor the would-be licensees perceive a need or desire to pursue either approach.¹⁸
- Reaffirm in any legislative or executive action recommendations that Congress and the Administration follow the strong copyright protection principles the U.S. agreed to in the Oct. 30, 2023 G7 Hiroshima Process International Guiding Principles for Organizations Developing Advanced AI System and Code of Conduct for Organizations Developing AI systems, as further described in the A2IM-RIAA Initial comments, and oppose any proposed text and data mining exceptions for AI.¹⁹

¹⁴ See, e.g., Isabela Raygoza, *Bad Bunny Slams 'Subpar' AI Soundalike That's Going Viral on TikTok*, Billboard (Nov. 8, 2023) ("If you happen to enjoy that subpar viral TikTok track, please leave this group immediately. You are not deserving of my friendship"), <https://www.billboard.com/music/latin/bad-bunny-slams-ai-soundalike-viral-tiktok-1235466825/>; Julia Gray, *All the Artists Who Have Spoken Out Against AI Music This Year*, The Messenger (Nov. 14, 2023 10:49 AM) (Artists speaking out include Bad Bunny, Sheryl Crow, Sting, John Legend, Ed Sheeran, and more) <https://themessenger.com/entertainment/artists-against-ai-music-quotes-bad-bunny>.

¹⁵ *Thaler v. Perlmutter*, __ F. Supp. 3d ___, 2023 WL 5333236, at *1 (D.D.C. Aug. 18, 2023).

¹⁶ Letter from Robert J. Kasunic, Associate Register of Copyrights and Director of Registration Policy and Practice, U.S. Copyright Office, to Van Lindberg, Re: Zarya of the Dawn 8-10 (Registration #Vau001480196), <https://www.copyright.gov/docs/zarya-of-the-dawn.pdf>.

¹⁷ Decision Re: Second Request for Reconsideration for Refusal to Register Théâtre D'opéra Spatial at 2 (SR #1-11743923581; Correspondence ID: 1-5T5320R), <https://copyright.gov/rulings-filings/review-board/docs/Theatre-Dopera-Spatial.pdf>.

¹⁸ See, e.g., A2IM-RIAA Initial Comments at 25-26; Andreessen Horowitz Initial Comments at 8; CCIA Initial Comments at 15; Library Copyright Alliance Initial Comments at 6; Internet Archive Initial Comments at 1, 12; Engine Initial Comments at 8; see also, n.178,179, *infra*.

¹⁹ Recently, Prof. Eleonora Rosati described why the text and data mining exceptions in Japan and Singapore should not allow unrestricted use of copyrighted content for the creating of training sets and AI training under the three-step test. See Eleonora Rosati, *No Step-Free Copyright Exceptions: The Role of the Three-step in Defining Permitted Uses of Protected content (including TDM for AI-Training Purposes)*, Nov. 10, 2023, SSRN: <https://ssrn.com/abstract=4629528> or <http://dx.doi.org/10.2139/ssrn.4629528>.

Discussion

Rather than reiterating the arguments and positions set forth in our Initial Comments, we incorporate those comments here by reference. Instead, we write now to respond to specific issues raised in others' Initial Comments that warrant a response.²⁰ In so doing, we limit our comments to topics that concern *generative* AI and we focus wherever possible on issues that arise uniquely in connection with sound recordings and music-related audio-visual works.²¹

I. Fair Use

We begin here because fair use is one of the most critical copyright issues that arises in connection with generative AI, and it is the one that has been mischaracterized the most by many of the AI Company Commenters.²²

*Fair use is not subject to categorical pronouncements.*²³ As the AI Company Commenters know (or should know), fair use is not subject to bright-line conclusions. Rather, it “requires judicial balancing, depending upon relevant circumstances”²⁴ and a thorough analysis of all four factors set forth in Section 107 of the Copyright Act.

²⁰ Given the vast number of comments filed and the short window of time for Reply Comments, we were not able to read or reply to all of the Initial Comments filed. Our failure to address any particular set of filed comments, therefore, should not be construed as agreement with or acquiescence in any such comments. To the extent that the Office is interested in our views on any issues that we do not address, either here or in our initial comments, we would be pleased to participate in an *ex parte* meeting to discuss them.

²¹ Most of the initial comments we reviewed in connection with this matter limit their analysis and discussion to large language models (“LLMs”). As the Office knows, not all generative AI systems are LLMs and what might be true with respect to LLMs may or may not be true of other types of generative AI models, including those designed to create new sound recordings, new audio-visual works, and new works of visual art.

²² See, e.g., CCIA Initial Comments at 5, 7, 12; Google Initial Comments at 1-2; Internet Archive Initial Comments at 8; Microsoft Initial Comments at 7-8; OpenAI Initial Comments at 11-12; Public Knowledge Initial Comments at 8-9, 11-12; Stability AI Initial Comments at 2, 8.

²³ See, e.g., CCIA Initial Comments at 5 (“any report issued by the Office as a result of this inquiry should clearly state that fair use permits the ingestion of copyrighted material in the course of an AI process”); Andreessen Horowitz Initial Comments at 7 (“Imposing infringement liability for the use of copyrighted works in AI model training, notwithstanding the case law that clearly demonstrates why such uses are fair, would be extremely misguided”); Anthropic Initial Comments at 9 (“Because training LLMs is a fair use, we do not believe that licensing is necessary *per se*”); Author’s Alliance Initial Comments at 14 (“[C]reators’ fears of job displacement . . . does not disturb the conclusion that the use of copyrighted works as training materials remains a fair use”); R Street Initial Comments at 4 (“The fair use doctrine should be interpreted in a manner that supports AI technology advancement”); Public Knowledge Initial Comments at 9 (“The AI training process itself does not violate any copyright rights, and need not even be considered through fair use analysis”); Engine Initial Comments at 5 (“For the sake of keeping the AI ecosystem innovative, competitive, and accessible to startups, it is most efficient to determine that the ingestion of copyrighted content as part of a training data set is a lawful, noninfringing use under copyright law, stopping any inquiries into infringement before the question of fair use arises”); Microsoft Initial Comments at 7 (“[T]he law is also clear that the use of copyrighted material is fair use”); Stability AI Initial Comments at 8 (We encourage policymakers to preserve this culture of ‘free learning’ in any future reform”).

²⁴ *Andy Warhol Foundation Visual Arts v. Goldsmith*, 598 U.S. 508, 527(2023) (quoting *Google LLC v. Oracle America, Inc.*, 141 S. Ct. 1183, 1197 (2021)). See also, e.g., *Thomson Reuters Enter. Ctr. GmbH v. Ross Intelligence Inc.*, No. 20-cv-613, 2023 WL 6210901, at *6-11 (D. Del. Sept. 25, 2023) (declining to grant summary judgment on fair use defense asserted by company that ingested copyrighted works into an AI system, finding instead that the issue must be decided by jury).

A fair use analysis must account for the entire context involved. Generative AI models are not developed in the abstract, and the purpose of feeding copyrighted content into them is not merely for the models to “learn” about those materials. Instead, generative AI models are often – and perhaps even typically – developed for the ultimate purpose of creating synthetic material that competes in the marketplace with the very material on which the model was “trained.” A proper fair use analysis of such copying cannot ignore that end purpose while myopically (and opportunistically) focusing on the curation and ingestion processes as ends unto themselves.²⁵ This is consistent with the *Google Books* decision cited by many of the AI Company Commenters and discussed further below, where it was Google’s limited end purpose – not to generate competing material, but rather to build a completely different product (a search engine) – that persuaded the court that the mass unauthorized copying was fair use, *albeit at its “boundaries.”*²⁶ Generative AI, on the other hand, far exceeds those boundaries. With generative AI models, the end purpose of producing competing content is the paradigm of an unfair use, as the Supreme Court made clear in *Warhol*.²⁷

Copying is not limited to the curation or ingestion stages. Proper consideration of the factual context of developing generative AI models must also take into account that generative AI models can and do embody expression copied from copyrighted inputs.²⁸ The embodiment of copyrighted inputs in a model may not be as obvious as in the case of more traditional copies, and while the copyrighted works represented in a generative AI model may be buried inscrutably among the numerical values making up the sometimes billions of nodes of a neural network, such embodiments nevertheless exist, similar to the numerical embodiments of sound recordings in the form of “zeros and ones” in compact discs or other digital formats. That copyrighted works can be and are embodied in generative AI models is a basic fact that has not been disputed by even ardent proponents of a sweeping application of fair use.²⁹

The search and plagiarism detection cases cited by AI Company Commenters are inapposite because no synthetic works that compete with the copyrighted works were created in those cases. When copyrighted works are used in developing generative AI models, the purpose is often to develop tools that generate synthetic outputs that are generally the same type of work as the inputs and are potential substitutes in the market for the inputs themselves. This is very different from the facts and circumstances involved in the cases on which the AI Company Commenters rely, where copying was done for purposes other than the creation of competing or substitutional works, and no new synthetic works were created:

²⁵ See *Andy Warhol Foundation for the Visual Arts, Inc. v. Goldsmith*, 598 U.S. 508, 533 (2023) (“The same copying may be fair when used for one purpose but not another”).

²⁶ *Id.* at 206; See also, *id.* at 216-18, 221-23, 224-25 (emphasis added).

²⁷ See *Warhol*, 598 U.S. at 528 (noting that copying for a substitutional purpose is “copyright’s bête noire”).

²⁸ We discuss these issues further in Section III, Ingestion, *infra* at 18-23.

²⁹ See, e.g., Anthropic Initial Comments at 3, 6-7, 11; CCIA Initial Comments at 9 n.9, 15-16, 20-21; Google Initial Comments at 13; Hugging Face Initial Comments at 7, 9; Microsoft Initial Comments at 8, 10-11; OpenAI Initial Comments at 7, 10, 14. See also Section III on Ingestion, *infra* at 18-23; Milad Nasr *et al.*, *Scalable Extraction of Training Data from (Production) Language Models*, arxiv.org, Nov. 28, 2023, <https://arxiv.org/pdf/2311.17035.pdf> (hereinafter “Nasr Article”).

- As described above, *Google Books* involved copying to serve the limited purpose of enabling searching.³⁰ The court found that purpose to be distinct from the purpose of the copied works.³¹ The court also credited the fact that Google took steps to protect the copies of books in its database, such as only showing short “snippets” of text to highlight a search term and implementing anti-hacking measures.³² Accordingly, the court determined that it was unlikely that Google’s copying could serve as a substitute for the copied works, finding that “Google’s making of a digital copy to provide a search function . . . augments public knowledge by making available information about [p]laintiffs’ books without providing the public with a substantial substitute for matter protected by the [p]laintiffs’ copyright interests in the original works or derivatives of them.”³³
- *iParadigms* involved archiving of student works uploaded to a plagiarism detection service under circumstances that led the district court to find the archiving to be licensed, a conclusion that the court of appeals did not address.³⁴ The archived papers were neither publicly displayed nor were they read or reviewed by iParadigms’ employees,³⁵ but rather were used to determine whether later-uploaded papers copied from them.³⁶ The court found that this purpose “was completely unrelated to expressive content” and so was transformative.³⁷
- *Kelly v. Arriba Soft* and *Perfect 10 v. Amazon* both involved reproduction of small, low-resolution “thumbnail” copies of images available on the internet and the display of those thumbnail images as search results.³⁸ The court was persuaded that the small size and low resolution of the thumbnails made them unsuited to the aesthetic purpose of the originals and, in effect, limited their use to indexing and improving access to images on the internet. This it found was “an entirely different function than” served by the originals.³⁹ The *Perfect 10* court simply followed *Kelly* in this respect.⁴⁰
- *Field v. Google*, a Nevada district court case, also involved internet searching, as well as Google’s caching of web pages for this purpose.⁴¹ The court found Google’s use to be impliedly licensed, and also held that the plaintiff was estopped from challenging the

³⁰ *Authors Guild v. Google*, 804 F.3d at 208-10.

³¹ *Id.* at 217-18.

³² *Id.* at 226-29.

³³ *Id.* at 207. Similarly, *HathiTrust* involved, in relevant part, copying to enable searching with *no* display of text to the searcher. *Authors Guild v. HathiTrust*, 755 F.3d 87, 91-92 (2d Cir. 2014). The court found the purpose of the copying to be transformative because, with no display of text to the searcher, the court could “discern little or no resemblance between the original text and the results of the HDL full-text search.” *Id.* at 97. As with *Google Books*, the copying in *HathiTrust* did not involve the creation of new works to compete with the copied material.

³⁴ *A.V. ex rel. Vanderhye v. iParadigms, LLC*, 562 F.3d 630, 634-35, 645 & n.8 (4th Cir. 2009).

³⁵ *Id.* at 641.

³⁶ *Id.* at 634.

³⁷ *Id.* at 640.

³⁸ *Perfect 10 v. Amazon.com, Inc.*, 508 F.3d 1146, 1155-56 (9th Cir. 2007); *Kelly*, 336 F.3d at 815.

³⁹ *Kelly*, 336 F.3d at 818-19.

⁴⁰ *See* 508 F.3d at 1165-67.

⁴¹ *Field v. Google Inc.*, 412 F. Supp. 2d 1106, 1114-15 (D. Nev. 2006).

use.⁴² As a third alternative ground for its decision (and thus in *dicta*), the court also found the use to be fair. In applying the first factor, the court noted that it was “clear that [Google] does not intend a ‘Cached’ link of a page to substitute for a visit [to] the original page.”⁴³ The court also gave weight to the fact that Field had opted in to search engine copying and caching his works by proactively “add[ing] the robots.txt file to his site to ensure that all search engines would include his Web site in the search listings.”⁴⁴

Because the purpose and character of the uses in these cases – all of which pre-date *Warhol* – are so different from those for generative AI, they are of no relevance in determining whether copying and ingesting copyrighted works for purposes of developing generative AI tools that create potentially competing synthetic works is fair use.

The exemption for text and data mining from the last triennial rulemaking was not a categorical blessing for AI-related copying. The Office’s analysis in the last triennial rulemaking of use of motion pictures and literary works for text and data mining is similarly misconstrued by the AI Company Commenters. There, the use permitted by the Office was subject to a number of explicit conditions, including a requirement that: all copies *be lawfully acquired*; the works be used solely for purposes of verifying research findings, not for their expressive purposes; and the institution storing the corpus of copyrighted works must implement “effective security measures” to safeguard the corpus from unauthorized access.⁴⁵ The Register found that the use involved was “likely transformative” because the purpose was to provide information about works in a manner different from “the expressive or informative purposes of the original works.”⁴⁶ Like the cases above, this rule pre-dated *Warhol* and the context is very different from generative AI development.

The intermediate copying cases cited by AI Company Commenters involved unprotected functional software code, not creative expression. The AI Company Commenters rely on two cases involving so-called “intermediate copying” of software to enable interoperability. As an initial matter, it is important to recognize that courts have not found “intermediate copying” to be a fair use *per se*.⁴⁷ Instead, they have evaluated so-called “intermediate copying” just like other copying – by careful application of the fair use factors to the particular facts and circumstances involved. In addition, both of these cases involved functional computer code, not highly creative

⁴² *Id.* at 1116-17.

⁴³ *Id.* at 1119.

⁴⁴ *Id.* at 1120.

⁴⁵ 37 C.F.R. § 201.40(b)(4), (5); 2021 Register’s Report, Eighth Triennial Proceeding, Recommendation of the Register of Copyrights at 102-03, 107-09 (Oct. 2021) https://cdn.loc.gov/copyright/1201/2021/2021_Section_1201_Registers_Recommendation.pdf.

⁴⁶ *Id.* at 109.

⁴⁷ *Fox Broadcasting Co. Inc. v. Dish Network, L.C.C.*, 905 F. Supp. 2d 1088, 1102-04 (C.D. Cal. 2012) (status of certain intermediate copies not resolved by *Sega* and must be analyzed under the four factors, *aff’d*, 747 F.3d 1060 (2014); *Sega Enterprises Ltd. v. Accolade*, 977 F.2d 1510, 1518 (9th Cir. 1992) (“the Copyright Act does not distinguish between unauthorized copies of a copyrighted work on the basis of what stage of the alleged infringer’s work the unauthorized copies represent”); *Walker v. University Books*, 602 F.2d 859, 864 (9th Cir.1979) (“[T]he fact that an allegedly infringing copy of a protected work may itself be only an inchoate representation of some final product to be marketed commercially does not in itself negate the possibility of infringement”); *Walt Disney Productions v. Filmation Associates*, 628 F.Supp. 871, 876 (C.D. Cal.1986) (“the Act prohibits the creation of copies, even if the creator considers those copies mere interim steps toward some final goal”).

expressive works like sound recordings, song lyrics, cover art, literary works, and the like, a distinction that is consequential in a fair use analysis.⁴⁸

- *Sega* involved copying of small portions of object code from a copyrighted game program “to gain an understanding of the unprotected functional elements of the program” in order for the defendant to write its own computer code for games that would be compatible with Sega’s game console.⁴⁹ The defendant employed a two-stage process to ensure that its games copied only unprotected functional elements of the plaintiff’s computer code. In finding that defendant’s reverse engineering for compatibility purposes constituted fair use, the court was clear that its analysis was specific to the functional computer code at issue. According to the court, because “Sega’s video game programs contain unprotected aspects that cannot be examined without copying, we afford them a lower degree of protection than more traditional literary works.”⁵⁰ In addition, the court considered the fact that defendant’s initial copying was the only way for it to access the unprotected elements of plaintiff’s work.⁵¹
- Similarly, *Connectix* involved copying a game console’s firmware to discern its unprotected functional elements as part of a reverse engineering process to enable creation of new software (a console emulator) that did not copy any protected expression from the original firmware.⁵² Noting the parallel to *Sega*, the court described the computer programs as “utilitarian articles,” subject to a lower degree of protection than purely expressive works, where the fair use doctrine “preserves public access to the ideas and functional elements embedded in” them.⁵³ Because it viewed the end console emulator as “modestly transformative,” and achieving software compatibility a legitimate purpose that supported competition, the court concluded that “the purpose and character of Connectix’s copying points toward fair use.”⁵⁴

Unlike the functional code at issue in *Sega* and *Connectix*, AI models use highly expressive works to create synthetic content that mimics the ingested works for purposes that go beyond mere interoperability. A copyrighted sound recording, music video, or cover art does not have functional aspects. Nobody needs to ingest such a work into an AI model to discern the ideas it expresses – one need only listen to it or look at it. Moreover, because AI’s outputs are not expressions of *human* creativity, they do not serve copyright’s purpose of promoting creative expression.⁵⁵ *Sega* and *Connectix* – both of which, like the other cases cited above, pre-date *Warhol* – have little relevance to the development of AI models.

⁴⁸ As the Supreme Court observed in *Google LLC v. Oracle Am., Inc.*, 141 S.Ct. 1183, 1202 (2021), functional computer code that is “inherently bound together with uncopyrightable ideas (general task division and organization)” is “further than are most computer programs ... from the core of copyright.”

⁴⁹ *Sega*, 977 F.2d at 1514; *see also id.* at 1522.

⁵⁰ *Id.* at 1526.

⁵¹ *Id.* at 1514, 1522-23.

⁵² *Sony Computer Entm’t, Inc. v. Connectix Corp.*, 203 F.3d 596, 598-601 (9th Cir. 2000).

⁵³ *Id.* at 603.

⁵⁴ *Id.* at 606-07.

⁵⁵ *See* Copyright Registration Guidance: Works Containing Material Generated by Artificial Intelligence, 88 Fed. Reg. 16190, 16192 (Mar. 16, 2023) [hereinafter “Copyright Registration Guidance”] (in the case of current generative AI technologies, the technology determines the elements of the outputs, and “the generated material is not the product of human authorship”).

AI Company Commenters generally fail to cite similar cases that found no fair use. There are a number of other cases that also involved search engines and reverse engineering that the AI Company Commenters fail to acknowledge in their Initial Comments. These include the following:

- *Associated Press v. Meltwater*⁵⁶ a Southern District of New York case where the court found that the challenged use – which involved scraping news articles from the web and providing excerpts from those articles to subscribers to its daily news clipping service – was not fair use, despite defendant’s attempts to characterize its service as a search engine that made transformative use of the copyrighted news articles. In rejecting Meltwater’s reliance on *Perfect 10* and *Kelly*, the court reasoned that “even if it were a search engine [which the court found it was not] it would still be necessary to examine whether Meltwater had acted to violate the Copyright Act... In other words, using the mechanics of search engines to scrape material from the Internet and provide it to consumers in response to their search requests does not immunize a defendant from the standards of conduct imposed by law through the Copyright Act, including the statutory embodiment of the fair use defense.”⁵⁷ The court also noted that “[i]nstead of driving subscribers to third-party websites, Meltwater News acts as a substitute for news sites operated or licensed by AP.”⁵⁸
- *Fox News Network v. TVEyes*,⁵⁹ a Second Circuit case involving a defendant that permitted users to search for and watch up to 10-minute clips of programming from over 1,400 TV and radio channels. The court noted that it had “cautioned [in *Google Books*] that the case ‘tested the boundaries of fair use.’”⁶⁰ It then found that the defendant “exceeded those bounds” and “failed to show that the product it offers to its clients can be justified as a fair use.”⁶¹
- *Disney v. VidAngel, Inc.*,⁶² a Central District of California case in which VidAngel made “intermediate copies” that it claimed were not “copies” as defined by the Copyright Act because the copies could not be viewed by consumers.⁶³ In addition to rejecting defendant’s fair use claim,⁶⁴ the court rejected VidAngel’s attempted reliance on *Sega v.*

⁵⁶ *Associated Press v. Meltwater*, 931 F. Supp. 2d 537 (S.D.N.Y. 2013).

⁵⁷ *Id.* at 556.

⁵⁸ *Id.* at 554.

⁵⁹ *Fox News Network, LLC v. TVEyes, Inc.*, 883 F. 3d 169 (2nd Cir. 2018).

⁶⁰ *Id.* at 174.

⁶¹ *Id.*

⁶² *Disney Enterprises, Inc. v. VidAngel, Inc.*, 224 F. Supp. 3d 957 (CD Cal. 2016), *aff’d on other grounds Disney Enterprises, Inc. v. VidAngel, Inc.*, 869 F. 3d 848 (9th Cir. 2017).

⁶³ The court described VidAngel’s process of copying as follows: “VidAngel creates ‘intermediate’ files’ by tag[ging] the files for over 80 types of content, and break[ing] them into approximately 1,300 fragments that contain no more than 10 seconds of content, then encrypt[ing] those fragments, and stor[ing] them in a secure, access-controlled location in the cloud.” *Id.* at 969.

⁶⁴ *Id.* at 974.

*Accolade*⁶⁵ and concluded that “VidAngel’s fragmented copies may not be able to be perceived directly by consumers, however they are able to be perceived with the aid of VidAngel’s software. Thus, the copying performed by Defendants falls within the category of acts that are proscribed by the statute.”⁶⁶

The above cases demonstrate that the mere fact that a party’s use involves some search function or intermediate copying does not itself end the fair use inquiry in favor of such use.

The AI Company Commenters also fail to cite to or appropriately take into account numerous cases where courts have declined to find that a use is transformative when the end product serves the same or similar purpose as the original,⁶⁷ or where the commercial nature of the use weighs against fair use.⁶⁸ These cases are closer to the generative AI context, where the output of the generative AI model often serves the same purpose as the original and the model often has some commercial purpose.

Because AI uses are highly commercial and barely transformative, if at all, factor 1 will rarely favor fair use. As explained by the Supreme Court in *Warhol*, to determine whether or not the first factor supports a finding of fair use, it is necessary to balance the degree to which a use may be considered transformative against the commercial or noncommercial nature of the use.⁶⁹

⁶⁵ In so doing, it observed that the *Sega* court had stated that “on its face, the language of 17 U.S.C. § 106(1) unambiguously encompasses and proscribes ‘intermediate copying.’” *Id.* at 1518. 17 U.S.C. § 101 provides that “in order to constitute a ‘copy’ for purposes of the Copyright Act, the allegedly infringing work must be fixed in some tangible form, “from which the work can be perceived, reproduced, or otherwise communicated, either directly or with the aid of a machine or device.”

⁶⁶ *Id.* at 970.

⁶⁷ See, e.g., *Warhol*, 598 U.S. at 533-36 (photograph of Prince and Warhol silkscreen of Prince served same purpose of illustrating magazine article, so licensing of latter to magazine was not transformative); *TCA Television Corp. v. McCollum*, 839 F.3d 168, 181-83 (2d Cir. 2016) (use of comedy sketch in play served same comedic purpose as original sketch); *Los Angeles News Serv. v. Reuters Television Int’l Ltd.*, 149 F.3d 987, 993-94 (9th Cir. 1998) (reproducing and airing news footage to report on the same news events was for the same intrinsic purpose as the original and not transformative); *Worldwide Church of God v. Philadelphia Church of God, Inc.*, 227 F.3d 1110, 1117 (9th Cir. 2000) (copying a religious book to create a new book for use by a different church audience was for the same intrinsic purpose as the original and not transformative).

⁶⁸ See, e.g., *Warhol* at 532 (“whether the use of a copyrighted work has a further purpose or different character, which is a matter of degree, and the degree of difference must be balanced against the commercial nature of the use”); *Harper & Row*, 471 U.S. at 562 (“The crux of the profit/nonprofit distinction is not whether the sole motive of the use is monetary gain but whether the user stands to profit from exploitation of the copyrighted material without paying the customary price”). See also *Hachette Book Grp., Inc. v. Internet Archive*, __ F. Supp. 3d ___, 2023 U.S. Dist. LEXIS 50749 (S.D.N.Y. 2023) (scanning of print books and digitally lending them without authorization is not fair use); *Worldwide Church of God v. Phila. Church of God, Inc.*, 227 F.3d 1110, 1118 (9th Cir. 2000) (copying by a nonprofit religious organization of an out-of-print book is not fair use); *Soc’y of Holy Transfiguration Monastery, Inc. v. Gregory*, 689 F.3d 29, 61 (1st Cir. 2012) (posting of plaintiff’s English translations of ancient religious texts on defendant Archbishop’s website did not constitute fair use); *Blackwell Publ’g, Inc. v. Excel Research Grp., LLC*, 661 F. Supp. 2d 786 (E.D. Mich. 2009) (copying by third party of copyrighted content in course packs without authorization to sell them to students was not fair use); *Am. Geophysical Union v. Texaco, Inc.*, 60 F.3d 913 (2d Cir. 1995) (copying of articles by researchers beyond what was permitted in the license for the articles was not fair use).

⁶⁹ *Warhol*, 598 U.S. at 525, 530-31, 532-33. The decision also considerably narrowed the relevance of a finding that a challenged use is “transformative.”

When copyrighted works are used to develop generative AI models, the two considerations will generally point in the same direction – against a finding of fair use.

Development of AI models is often a commercial activity that is being pursued by for-profit businesses, including the large technology companies that provided initial comments, as a major focus of business. As one large investor in AI companies explained, AI development has attracted “billions and billions of dollars” in investments.⁷⁰ Even where AI development is ostensibly being carried out by nonprofit organizations,⁷¹ there may be contractual arrangements that make the development commercial in effect. For example, Open AI has a non-profit parent company, but carries out most of its work through a for-profit affiliate.⁷² It has received billions of dollars in investments and exclusively licensed significant results of its development efforts to Microsoft.⁷³ And while Anthropic is apparently organized as a public benefit corporation, there are reports that Google has agreed to invest up to \$2 billion in Anthropic⁷⁴ and Amazon has reportedly invested up to \$4 billion.⁷⁵ To be sure, there is academic research involving AI development as well, but “[f]unding from for-profit developers may blur the noncommercial nature of research projects.”⁷⁶

The facts of each particular case will matter, but AI development often is not only commercial, but big business. Where development of AI models is commercial, that “tends to weigh against a finding of fair use.”⁷⁷

Factor 2 weighs against fair use when creative works are ingested. The second fair use factor is “the nature of the copyrighted work.”⁷⁸ Application of this factor will obviously depend on the particular types of works used by the defendant, because “some works are closer to the core of intended copyright protection than others.”⁷⁹ The second factor will clearly weigh against fair

⁷⁰ Andreessen Horowitz Initial Comments at 6.

⁷¹ See A2IM-RIAA Initial Comments at 19 and n.74; see also *American Geophysical Union v. Texaco, Inc.*, 60 F. 3d 913, 924 (2d Cir. 1994) (“The purposes illustrated by the categories listed in section 107 refer primarily to the work of authorship alleged to be a fair use, not to the activity in which the alleged infringer is engaged. Texaco cannot gain fair use insulation for [its employee]’s archival photocopying of articles (or books) simply because such copying is done by a company doing research.”),

⁷² OpenAI LP (March 11, 2019), <https://openai.com/blog/openai-lp>.

⁷³ Kyle Wiggers, *Microsoft invests billions more dollars in OpenAI, extends partnership*, TechCrunch (Jan. 23, 2023), <https://techcrunch.com/2023/01/23/microsoft-invests-billions-more-dollars-in-openai-extends-partnership/?guccounter=1>; OpenAI, *OpenAI licenses GPT-3 technology to Microsoft* (Sept. 22, 2020), <https://openai.com/blog/openai-licenses-gpt-3-technology-to-microsoft>; Microsoft, *OpenAI forms exclusive computing partnership with Microsoft to build new Azure AI supercomputing technologies* (July 22, 2019), <https://news.microsoft.com/2019/07/22/openai-forms-exclusive-computing-partnership-with-microsoft-to-build-new-azure-ai-supercomputing-technologies/>.

⁷⁴ CNBC, *Google commits to invest \$2 billion in OpenAI competitor Anthropic* (Oct. 27, 2023), <https://www.cnbc.com/2023/10/27/google-commits-to-invest-2-billion-in-openai-competitor-anthropic.html>.

⁷⁵ See <https://techcrunch.com/2023/09/25/amazon-to-invest-up-to-4-billion-in-ai-startup-anthropic/>.

⁷⁶ Public Knowledge Initial Comments at 12.

⁷⁷ *Harper & Row v. Nation Enters.*, 471 U.S. 539, 562 (1985) (quoted in *Warhol*, 598 U.S. at 537).

⁷⁸ 17 U.S.C. § 107(2).

⁷⁹ *Campbell v. Acuff-Rose Music, Inc.*, 510 U.S. 569, 586 (1994).

use when an AI developer uses highly creative works like sound recordings, cover art, music videos, or song lyrics, which are “close to the core of the intended copyright protection.”⁸⁰

Factor 3 will not favor fair use because AI systems typically ingest whole works. The third fair use factor is “the amount and substantiality of the portion used in relation to the copyrighted work as a whole.”⁸¹ “[A] finding of fair use is more likely when small amounts, or less important passages, are copied than when the copying is extensive, or encompasses the most important parts of the original.”⁸² Copyrighted works used in development of generative AI models are typically used in their entirety, even if they are broken up into pieces (or tokens) as part of the AI development process. It is well settled that “copying an entire work militates against a finding of fair use.”⁸³

The AI Company Commenters advocate that the use of any particular work is negligible because AI models ingest so many works that the influence of any one of them on the overall model is limited.⁸⁴ However, that argument has the third factor backward. As the text of Section 107(3) makes clear, the third factor is about how much of the *underlying* work is used, not how big a part it is of any new work. Courts agree that the factor is to be analyzed “with reference to the copyrighted work, not the infringing work.”⁸⁵

Factor 4 weighs against fair use because use of copyrighted works in generative AI models can harm the markets for those works. The fourth factor is “the effect of the use upon the potential market for or value of the copyrighted work.”⁸⁶ As explained in the A2IM-RIAA Initial Comments, the inquiry here is whether, “if the challenged use ‘should become widespread, it would adversely affect the potential market for the copyrighted work.’”⁸⁷ When copyrighted works are used in the development of a generative AI model that outputs the same type of works, the potential for those outputs to supplant the market for the input works is clear.⁸⁸

The *Sega* and *Connectix* cases are not to the contrary, as those cases involved copying of unprotected functional elements of software (not highly expressive works like sound recordings)

⁸⁰ *Id.*; see also *UMG Recordings, Inc. v. MP3.com, Inc.*, 92 F. Supp. 2d 349, 351 (S.D.N.Y.2000); see also *A&M Records, Inc. v. Napster, Inc.*, 239 F.3d 1004, 1016 (2001); *Capitol Records, LLC v. ReDigi Inc.*, 934 F.Supp.2d 640, 654 (2013).

⁸¹ 17 U.S.C. § 107(3).

⁸² *Authors Guild v. Google*, 804 F.3d at 221.

⁸³ *Worldwide Church of God*, 227 F.3d at 1118 (internal quotation marks omitted).

⁸⁴ See, e.g., Library Copyright Alliance Initial Comments at 1; Meta Initial Comments at 15; Consumer Technology Association Initial Comments at 7; Engine Initial Comments at 8; New Media Rights Initial Comments at 17; Wikimedia Foundation Initial Comments at 8.

⁸⁵ *Bill Graham Archives v. Dorling Kindersley Ltd.*, 448 F. 3d 605, 613 (2d Cir. 2006); see also, e.g., *Google v. Oracle*, 141 S. Ct. at 1204-06 (analyzing number of lines of Java code copied by Google as a proportion of the original, not as a proportion of Google’s software incorporating them); *Dr. Seuss Enterprises, LP v. Comicmix LLC*, 983 F. 3d 443, 456-58 (9th Cir. 2020) (analyzing quantity and quality of copied material as a proportion of the original, not as a proportion of the work incorporating them).

⁸⁶ 17 U.S.C. § 107(4).

⁸⁷ *Harper & Row*, 471 U.S. at 568 (1985) (quoting *Sony*, 464 U.S. at 451).

⁸⁸ *Warhol*, 598 U.S. at 528 (“The ‘central’ question it [the first factor] asks is ‘whether the new work merely ‘supersede[s] the objects’ of the original creation ... (‘supplanting’ the original), or instead adds something new, with a further purpose or different character’”)(citations omitted).

and the court regarded the market harm from the copying in those cases as negligible.⁸⁹ Likewise, in *Google Books* and *HathiTrust*, the courts found that the fourth factor did not weigh against a finding of fair use because the identified harm was speculative and trivial.⁹⁰ Those cases are quite different from the circumstances likely to be presented in cases involving use of copyrighted works in development of generative AI models whose outputs can substitute for the copied works.

Today, content licensing is at the core of record companies' businesses. From a record company perspective, generative AI simply represents a new potential market for licensing uses of their sound recordings, similar to the way record companies have – for decades – been licensing their catalogues to a variety of innovative digital services, as well as derivative uses of sound recordings for use in audio-visual works (e.g., synchronization licensing) and licensed samples of recordings for use in new recordings. Record companies also license major social media platforms to offer full catalog libraries of music to consumers for use by those consumers in the creation of user-generated video content. And now there is a rapidly emerging market for licensing copyrighted works for use in developing generative AI models. This is true both within the recorded music industry⁹¹ and elsewhere.⁹² Indeed, even while denying that licensing is necessary for the use of copyrighted works in developing AI models, OpenAI acknowledges that it is “actively engaged in partnership discussions with copyright owners” and has “been told that many other leading AI companies are also actively pursuing similar partnerships.”⁹³

Taken together, the four factors weigh strongly against fair use. After the fair use factors are evaluated individually, they are to be “weighed together, in light of the purposes of copyright.”⁹⁴ Because the analysis of each of the factors will depend on the facts of the case involved, so too will their weighing. However, the foregoing analysis makes clear that, in most cases in which copyrighted works are used in developing generative AI models, all the factors are likely to weigh *against* fair use. Moreover, the purpose of copyright is to promote human creative endeavors.⁹⁵ That purpose is served by protecting human creators from having their works used to develop generative AI models that threaten to displace human creators by producing outputs

⁸⁹ *Connectix*, 203 F.3d at 607-08 (fourth factor weighs in defendant's favor notwithstanding the possibility of “some economic loss”); *Sega*, 977 F.2d at 1523-24 (fourth factor weighs in defendant's favor notwithstanding the possibility of “minor economic loss”).

⁹⁰ *Authors Guild v. Google*, 804 F.3d at 223-25; *Authors Guild v. HathiTrust*, 755 F.3d at 99-101.

⁹¹ See, e.g., Stability AI Initial Comments at 5 (in developing Stable Audio, StabilityAI used “800,000 soundtracks, obtained through a data access agreement with AudioSparx, a leading content library”). See, *supra*, at 2; see also A2IM-RIAA Initial Comments at 3, 13.

⁹² See, e.g., SIIA Initial Comments at 2 (“Many of our members already license their works for use as AI training data”); Open AI Initial Comments at 5 (To develop its models, OpenAI uses “information that we obtain from third parties through commercial arrangements” as well as internally-generated material and material found on the internet.); Adobe Initial Comments at 2 (In developing its first Firefly model, Adobe used “licensed images from [its] own Adobe Stock photography collection, openly licensed content, and public domain images”).

⁹³ OpenAI Initial Comments at 13.

⁹⁴ *Campbell*, 510 U.S. at 578.

⁹⁵ See Copyright Registration Guidance, 88 Fed. Reg. at 16191-92.

that do not embody human creativity while supplanting works of human creativity in the marketplace. Such uses will rarely, if ever, be fair uses.⁹⁶

II. The Need for a Federal “Voice and Likeness” Right

The initial comments provide broad support for a uniform federal right. Existing state right of publicity laws provide some protection against misappropriation of an individual’s identity in AI-generated outputs and, potentially, even use in training.⁹⁷ However, the scope of that protection varies from state to state, leaving individuals vulnerable to abuses of their voices and likenesses in AI-generated content depending on the jurisdiction in which they live. There appears to be a consensus among many of the commenters – including the AI Company Commenters – that a new federal right of individuals to control the use of their voices and likenesses is necessary to address the unique threats of generative AI.⁹⁸ Indeed, as YouTube recently recognized in announcing its new policies regarding AI-generated content uploaded to its platform, the unique risks of generative AI require new measures to protect against the harms presented “where someone’s face or voice could be digitally generated without their permission or to misrepresent their points of view.”⁹⁹

While general agreement in favor of a new federal right appears to be broad-based, commenters disagree on several aspects of what this new right should involve, including: (1) the manner in which First Amendment defenses to unauthorized uses of voices and likeness should be reflected in any new law, (2) the scope of secondary liability for violations of the right, and (3) whether the new law should preempt existing state rights of publicity. We address each of these points in turn.

The First Amendment provides a defense, but not a categorical one. It is well settled that the First Amendment provides a defense for unauthorized uses of an individual’s voice, image, or

⁹⁶ Some initial comments cite *Eldred v. Ashcroft*, 537 U.S. 186 (2003) as part of their discussion of fair use. However, that decision concerned the power of Congress to extend the duration of copyright protection. It merely noted the existence of the fair use doctrine and had nothing to say about how to apply it. Some initial comments also cite *Sony Corp. of Am. v. Universal City Studios* as part of their discussion of fair use. While that decision did involve a fair use determination with respect to the propriety of certain “time shifting,” 464 U.S. at 447-55, the AI Company Commenters focus on the decision’s discussion of contributory infringement, *id.* at 434-42. That discussion has no bearing on a fair use analysis. Instead, we discuss the *Betamax* case, *infra*, in connection with secondary liability.

⁹⁷ See Rothman Initial Comments at 3-4.

⁹⁸ See, e.g., Adobe Initial Comments at 8 (“Adobe has proposed that Congress establish a new Federal Anti-Impersonation Right (the “FAIR” Act) to address this type of economic harm”); BSA Initial Comments at 11 (“it is worth exploring how federal legislation could work to afford athletes, musicians, singers, actors, and other artists, performers, and public figures with the opportunity to prevent others from unauthorized exploitation of their name, image, likeness, or voice”); Microsoft Initial Comments at 12-13 (“given the capabilities of AI to generate convincing replicas, there is a potential need to provide clarity and certainty at the federal level for protection of digital replicas—unauthorized depictions of names, images, likenesses, and signatures, created without the consent of the depicted artist and used in commerce”); Public Knowledge Initial Comments at 4, 20 (“We broadly support the creation of an individual right against being digitally replicated via GAI”); International Trademark Association “INTA”) Initial Comments at 9-11 (“Congress should establish a new federal right of publicity with basic principles that would apply regardless of the means used to generate potentially infringing content”).

⁹⁹ See Jennifer Flannery O'Connor and Emily Moxley, *Our Approach To Responsible AI Innovation*, YouTube Official Blog (Nov. 14, 2023) <https://blog.youtube/inside-youtube/our-approach-to-responsible-ai-innovation/>.

likeness in certain circumstances. And as businesses whose core assets are creative works of human expression, we regard the First Amendment as an essential part of who we are and a bedrock principle of what we do. But we disagree with those commenters who argue that any federal right of an individual to control uses of their voice or likeness must contain express, categorical exclusions for all uses of a certain type, such as unauthorized uses of an individual's voice or likeness in any "expressive works," regardless of the particular facts and circumstances of the use.¹⁰⁰ The First Amendment requires no such thing; to the contrary, the inclusion of rigid exclusions for entire categories of uses risks leaving individuals unprotected against misappropriation of their voices and likenesses that the First Amendment would not protect.

In fact, existing right of publicity laws have long been applied to expressive works, and numerous courts have found that the right of publicity in such cases outweighs First Amendment interests when balanced on the facts of the particular use at issue. For example, the United States Supreme Court has recognized that video games are "expressive works" entitled to the same First Amendment protections as "the protected books, plays, and movies that preceded them"¹⁰¹ Yet, contrary to some commenters' assertions, we know of no court decision that has ever recognized a categorical First Amendment exemption for video games against right of publicity claims. Instead, numerous courts – applying a fact-based balancing analysis that weighs the right of publicity against the First Amendment interests at stake – found right of publicity claims to outweigh the First Amendment interests in unauthorized uses of individual likenesses in these "expressive works."¹⁰² Indeed, just this term, the Supreme Court held – in the analogous context of First Amendment defenses asserted against trademark claims – that "[t]here is no threshold test working to kick out all cases involving 'expressive works.'"¹⁰³

Simply put, there is no constitutional or jurisprudential basis for the unfounded contention that the First Amendment requires categorical exclusions from right of publicity claims against unauthorized uses of an individual's name, voice, or likeness. While some (but not all) state right of publicity laws include categorical exclusions for certain uses (and those exclusions are not consistent from state to state), those exclusions are the result of lobbying negotiations that were unique to the political circumstances of each state and are not compelled by either the First Amendment or judicial precedent. Rather, as courts have repeatedly found, the First Amendment calls for a case-specific balancing of the right of publicity against whatever First Amendment

¹⁰⁰ See MPA Initial Comments at 71 ("[T]he right of publicity does not—and, to be consistent with the First Amendment, may not—regulate uses of or references to individuals' NIL in 'expressive works,' such as books, plays, news articles and broadcasts, songs, and movies and television programs").

¹⁰¹ *Brown v. Ent'mt Merchants Ass'n*, 564 U.S. 786 (2011).

¹⁰² See, e.g., *Keller v. Electronic Arts, Inc. (In re NCAA Student-Athlete Name & Likeness Licensing Litigation)*, 724 F.3d 1268 (9th Cir. 2013) (rejecting video game maker's First Amendment defense against right of publicity claims by former college athletes whose likenesses were used in a video game without permission); *Hart v. Electronic Arts, Inc.*, 717 F.3d 141 (3d Cir. 2013) (acknowledging that "video games are protected as expressive speech under the First Amendment" but nevertheless rejecting First Amendment defense against right of publicity claims by former college athletes whose likenesses were used in a video game without permission); *Davis v. Electronic Arts, Inc.*, 775 F.3d 1172 (9th Cir. 2015) (rejecting First Amendment defense against right of publicity claims by former professional football players whose likenesses were used in a video game without permission); *No Doubt v. Activision Publishing, Inc.*, 192 Cal. App.4th 1018 (2011) (rejecting First Amendment defense against right of publicity claims by well-known rock band whose likenesses were used in a video game without permission).

¹⁰³ *Jack Daniel's Properties, Inc. v. VIP Prods. LLC*, 143 S. Ct. 1578, 1591 (2023).

interests may be presented in the given case. It is precisely this sort of balancing approach that should be reflected in any federal right.

The only Supreme Court case to date that has ever examined the right of publicity – *Zacchini v. Scripps-Howard Broadcasting Co.*, 433 U.S. 562 (1977) – illustrates the danger of applying rigid, categorical exclusions from the scope of an individual’s right to control the use of their likeness. There, the Supreme Court balanced the plaintiff’s right of publicity claim against the First Amendment interests of a news station that had broadcast the plaintiff’s performance as part of a news story, without the performer’s permission. The Court found that, on the facts of that case, the plaintiff’s interest outweighed the First Amendment interests of the news station and allowed the claim to proceed. If the applicable right of publicity law had included a categorical exclusion for “news articles and broadcasts” – as at least one commenter wrongly suggests is constitutionally required – then a claim that the Supreme Court concluded was meritorious would not have been allowed to proceed at all.

There is no basis for a secondary liability carve-out. Several of the AI Company Commenters expressed concern about the potential scope of secondary liability for violations of a new federal right and argue that “any federal right of publicity would have to ensure that it does not inadvertently impose liability on platforms merely for hosting user-provided content.”¹⁰⁴ This concern seems misplaced, as secondary liability has long been a feature of existing state law regarding the right of publicity, so there is nothing new to the idea that a federal law should provide the same.¹⁰⁵ Indeed, secondary liability is a longstanding principle of tort law more generally, including in the context of intellectual property claims.¹⁰⁶ There is no reason why a federal right of individuals to control the uses of their voices and likenesses should not also include secondary liability for those who knowingly facilitate violations of the right by third parties, including by hosting AI-generated content with knowledge that the use of an individual’s voice or likeness in that content is unauthorized.

The concern of the AI Company Commenters most likely relates to the current circuit split over whether Section 230 of the Communications Decency Act¹⁰⁷ immunizes hosting platforms from liability for violations of state rights of publicity or whether such claims fall within the exception for intellectual property laws.¹⁰⁸ As one commenter noted, “[t]his matters because if the exception does not apply, it is difficult to get platforms to take down infringing content, and it can be hard to track down individuals who initially created or circulated works.”¹⁰⁹ Although the

¹⁰⁴ Google Initial Comments at 16.

¹⁰⁵ See, e.g., *Perfect 10, Inc. v. Cybernet Ventures, Inc.*, 213 F. Supp.2d 1146 (C.D. Cal. 2002) (holding that a defendant may be liable for “aiding and abetting” a right of publicity violation).

¹⁰⁶ See *id.* at 1183-84 (citing the Restatement (Second) of Torts §876). See also, e.g., *Gershwin Pub. Corp. v. Columbia Artists Management, Inc.*, 443 F.2d 1159, 1162 (2d Cir. 1971) (deriving secondary copyright infringement from general tort principles); *Tiffany (NJ) Inc. v. eBay Inc.*, 600 F.3d 93, 103-104 (2d Cir. 2010) (same as to secondary trademark infringement) (citations omitted).

¹⁰⁷ 47 U.S.C. § 230.

¹⁰⁸ Compare, e.g., *Perfect 10 v. CCBill, Inc.*, 488 F.3d 1102, 1107–08 (9th Cir. 2007) (holding that Section 230’s intellectual property exemption does not apply to state law claims) with, e.g., *Hepp v. Facebook*, 14 F.4th 204, 209–12 (3d Cir. 2021) (disagreeing with the Ninth Circuit and holding that both state and federal IP claims are carved out of Section 230); *Atlantic Recording Corp. v. Project Playlist*, 603 F. Supp.2d 690 (S.D.N.Y. 2009) (holding that “Section 230(c)(1) does not provide immunity for either federal or state intellectual property claims”).

¹⁰⁹ Rothman Initial Comments at 4-5.

Ninth Circuit stands alone among U.S. courts in holding that Section 230 immunizes platforms from state rights of publicity claims (and has been widely criticized for doing so),¹¹⁰ the major platforms are based in that Circuit and therefore, by sheer virtue of geography, enjoy an unfair – and unintended – immunity from those claims, even where they have knowledge that the use of an individual’s likeness is unauthorized. It is no wonder, then, that some of these platforms argue for maintaining the status quo.¹¹¹

To address this unfairness and allow for meaningful enforcement of voice and likeness rights, it is imperative that any law establishing a federal “voice and likeness” right make clear that it is an *intellectual property* right, which, as a federal law, would be carved out of Section 230’s immunity, even under the Ninth Circuit’s misreading of Section 230.

Any new federal right should not preempt state law. Several commenters suggest that any new federal “voice and likeness” right should preempt existing state right of publicity laws.¹¹² As even one of these commenters herself acknowledges, however, the new law “should not destabilize state right of publicity and privacy laws in this area which have been operating for more than 100 years.”¹¹³ It is precisely for this reason that a federal right should *not* preempt state law to deprive individuals of rights that have been carefully developed over decades of legislation and litigation. Moreover, the Senate’s discussion draft of the “NO FAKES Act” addresses only misuses of voice and likeness in connection with “digital replicas,” and thus addresses far fewer uses of names, images, voices, and likeness than are covered under current state laws.¹¹⁴ To preempt those state laws would risk depriving individuals of longstanding protections against unauthorized uses of their identities outside of the generative AI context. Rather than preempt existing state law, a new federal right should provide at least a minimum “floor” of protection for individuals from misuses of their voices and likenesses in the generative AI content.

III. Ingestion Involves Unauthorized Copying of Copyrighted Expression

The AI Company Commenters paint a misleading picture of the copyright implications of use of copyrighted works in the development of generative AI models. Contrary to their claims, development of generative AI models involves reproduction of copyrighted (as well as non-copyrighted) works and computer processing of expressive aspects of those works. In this section we address a number of the misstatements and misleading implications that some AI Company Commenters make.

Publicly available works are not free for the taking. The AI Company Commenters wrongfully claim that just because copyrighted works have been made “publicly available,” AI developers

¹¹⁰ See, e.g., *Project Playlist*, 603 F. Supp.2d at 703 (criticizing Ninth Circuit’s position as “lack[ing] any support in the language of the CDA”).

¹¹¹ See, e.g., Meta Initial Comments at 21 (“There is no reason to believe that legal doctrines like ... the right of publicity are inadequate to the task...”).

¹¹² See, e.g., EFF Initial Comments at 7; Rothman Initial Comments at 4.

¹¹³ Rothman Initial Comments at 4.

¹¹⁴ https://www.coons.senate.gov/imo/media/doc/no_fakes_act_draft_text.pdf.

have the right to take them and use them without permission or compensation.¹¹⁵ As the Office knows (and the AI Company Commenters should know), just because something is publicly available or accessible does not mean that it lacks copyright protection or that such protection was somehow waived. It simply means that a copyrighted work has been published.¹¹⁶ Ownership of (or access to) a copy or phonorecord and ownership of the copyright are two distinct issues.¹¹⁷ Just as one who owns (or accesses) a CD does not own the copyright in either the musical works or sound recordings embodied on the CD, one who is able to access a copyrighted work made available online, such as sound recordings available for streaming via various services, does not own the copyright in such work or have the unrestricted right to use the work in a way that violates the exclusive rights of the copyright owner under Section 106 of the Copyright Act.

Datasets include copies or phonorecords subject to the exclusive rights of copyright owners.

Some of the AI Company Commenters suggest that the process of ingesting copyrighted works into generative AI models involves only transitory copying, while ignoring the datasets so ingested.¹¹⁸ These datasets “are not merely a grab-bag of copyrighted works, but rather a highly curated set of data.”¹¹⁹ That curation is not happening on the fly. For example, the Common Crawl data set “contains petabytes of data, regularly collected since 2008” and “stored on Amazon Web Services’ Public Data Sets” and available for download.¹²⁰ These datasets clearly implicate the exclusive rights of copyright owners under Section 106.¹²¹ This is true even if copyrighted works are chopped up into tokens;¹²² those tokens still reproduce the works involved. If datasets are distributed to third parties, doing so implicates copyright owners’ exclusive distribution right under Section 106(3).¹²³

Copies made in connection with AI development are sufficiently persistent to constitute copyright infringement. The copying done in connection with AI development is not of “transitory

¹¹⁵ See, e.g., Business Software Alliance Initial Comments at 8; Microsoft Initial Comments at 6, 9; Google Initial Comments at 11; Ad Hoc Group of Developers & Users of Generative AI Initial Comments at 1; Public Knowledge Initial Comments at 9; Hugging Face Initial Comments at 1,6; Anthropic Initial Comments at 5; Chamber of Progress at 5; Duolingo Initial Comments at 2; Andreessen Horowitz Initial Comments at 10.

¹¹⁶ See 17 U.S.C. § 101 (definition of publication).

¹¹⁷ See 17 U.S.C. § 202 (ownership of copyright as distinct from ownership of material object).

¹¹⁸ For descriptions of datasets, see, e.g., Microsoft Initial Comments at 6 (referring to the Common Crawl and The Pile); Samuelson Initial Comments at 21 (not just collections of links); CCIA Initial Comments at 10 (describing sizes of various data sets).

¹¹⁹ TechNet Initial Comments at 11 (among other things, “resources are devoted to filtering data to (for example) remove duplicates and scrub potentially sensitive information, and to pre-processing data into a format suitable for model training”).

¹²⁰ Common Crawl, Overview, <https://commoncrawl.org/overview>.

¹²¹ See 17 U.S.C. § 101 (definitions of copies and phonorecords).

¹²² See, e.g., OpenAI Initial Comments at 6-7; BSA Initial Comments at 2; Microsoft Initial Comments at 6-7; Meta Initial Comments at 3-4; Stability AI Initial Comments at 10; Google Initial Comments at 4; Allen Institute for AI Initial Comments at 1, 7, 9; Hugging Face Initial Comments at 9; Workshop on Generative AI and Law Initial Comments at 37; Samuelson *et al.* Initial Comments at 8, 20.

¹²³ SIAA Initial Comments at 3 (“firms that are selling copies of others’ copyrighted works for training purposes face a ... liability risk”).

duration,” as a number of the AI Company Commenters claim.¹²⁴ In fact, the copies may be made at various stages of the AI development process¹²⁵ and are often retained for extended periods of time.¹²⁶ The two leading cases that address questions of fixation and duration are *MAI v. Peak*¹²⁷ and *Cartoon Network v. CSC Holdings*¹²⁸ (i.e., the *Cablevision* case). In *MAI*, the Ninth Circuit found that “the representation [of operating system software] created in the RAM [random access memory] is ‘sufficiently permanent or stable to permit it to be perceived, reproduced, or otherwise communicated for a period of more than transitory duration.’”¹²⁹ It further concluded that the “loading of MAI’s operating software into RAM, which occurs when an MAI system is turned on, constitutes a copyright violation.”¹³⁰ In *Cablevision*, the Second Circuit distinguished *MAI v. Peak*, finding that when audiovisual content was buffered in a computer server for 1.2 seconds solely as part of the process of passing that content on to another server, and was automatically erased and replaced every tenth of a second, the buffered embodiment persisted for only a transitory period, and thus was not fixed.¹³¹ It is our understanding that AI tools ingest copyrighted works into computer memory for purposes that are sufficiently permanent or stable to meet both the embodiment test set forth in *MAI* and are not so transitory as to implicate the 1.2 second threshold set forth in *Cablevision*.¹³²

The ingestion process results in the embodiment of copyrighted expression in generative AI models. The AI Company Commenters try to negate liability for unauthorized reproductions of copyrighted works by anthropomorphizing the AI development process and making it sound as if computers are “learning” information from input works the way that humans do. However, ingesting works into a generative AI model is not “like the act of reading a book,” and computers do not develop “knowledge.”¹³³ Computers are inanimate machines that are not capable of perceiving ideas, concepts, or even facts in the creative works they process.¹³⁴ Instead, the

¹²⁴ See, e.g., BSA Initial Comments at 8 (“the fleeting nature of the RAM copies involved in the machine learning process are unlikely to meet the Copyright Act’s threshold for ‘fixation’ and are therefore unlikely to be considered ‘copies’ that implicate the reproduction right”) (citations omitted).

¹²⁵ For a description of the different stages of the AI development process, see, e.g., Google Initial Comments at 5; BigBear.ai Initial Comments at 24-26; Adobe Initial Comments at 3; OpenAI Initial Comments at 6-7; BSA Initial Comments at 5-6; Katherine Lee *et al.* Initial Comments; Samuelson, *et al.*, Initial Comments at 19-20.

¹²⁶ See, e.g., Hugging Face Initial Comments at 6; BigBear.ai Initial Comments at 9-10; SeaQVN Initial Comments at 22-24; Public Knowledge Initial Comments at 6-7.

¹²⁷ *MAI Systems Corp. v. Peak Computer, Inc.* 991 F.2d 511,518 (9th Cir. 1993).

¹²⁸ *Cartoon Network v. CSC Holdings*, 536 F.3d 121, 124-25, 127-30 (2d Cir. 2008).

¹²⁹ *MAI Systems Corp. v. Peak Computer, Inc.* 991 F.2d 511,518 (9th Cir. 1993).

¹³⁰ *Id.*

¹³¹ *Cartoon Network v. CSC Holdings*, 536 F.3d 121, 124-25, 127-30 (2d Cir. 2008).

¹³² In addition, the intensive processing that occurs as part of the ingestion process is a far cry from simple buffering at issue in *Cablevision*, where the material involved was embodied only for purposes of passing it on to another server.

¹³³ Google Initial Comments at 9. Google’s comments claim that the Court in *Harper & Row* found “knowledge harvesting” to further the purposes of copyright law. But that is not what the Court said. To the contrary, the Court found that “the Second Circuit gave insufficient deference to the scheme established by the Copyright Act for fostering the original works that provide the seed and substance of this harvest. The rights conferred by copyright are designed to assure contributors to the store of knowledge a fair return for their labors.” *Harper & Row v. Nation Enters.*, 471 U.S. 539, 545-46 (1985).

¹³⁴ See Anthropic Initial Comments at 3. See also EFF Initial Comments at 2-3 (referring to “learning” and “plot twists”); Microsoft Initial Comments at 8 (analogizing to a person analyzing and learning from a work); OpenAI Initial Comments at 7 (analogizes to person who reads book and sets it down, but omits all the copying involved in the AI development process).

ingestion process starts with copyrighted human expression, such as the text of a literary work or a digital copy of a sound recording or image, and then computers perform calculations on the expressive aspects of the copyrighted works to represent that expression numerically.¹³⁵ This is merely another (albeit more complex) form of reproducing the copyrighted work.¹³⁶

Some AI Company Commenters, like OpenAI, argue that “[w]hen undergoing pre-training, a model is not interested in the expressive aspects of individual copyrighted works.”¹³⁷ Recent research shows that this is false.¹³⁸ Moreover, contrary to the AI Company Commenters’ claim that AI models are “not interested in the expressive aspects individual copyrighted works,” it is precisely the expressive content of the ingested copyrighted works (in our case sound recordings and related audiovisual works) that has value when a generative AI model is used to generate music or clone a voice.¹³⁹

Ingestion and processing of copyrighted works clearly implicates the reproduction right. Even though input works are represented in a massive matrix of numbers in a particular AI model, the fact that the input works can subsequently be reproduced in their entirety (or even in substantial part) as an AI output satisfies the statutory definitions of copy and phonorecord (as applicable). Those embodiments are defined in the Copyright Act as “material objects” in which a work is “fixed by any method now known or later developed” and from which it “can be perceived, reproduced, or otherwise communicated, either directly or with the aid of a machine or device.” Here, despite having been dissected and manipulated through the ingestion and model-building process, works nevertheless are sometimes capable of being further “perceived” or “reproduced” at the output stage with the aid of the AI “machine or device,” which merely confirms the fact that the works have been reproduced and are embodied in the AI model itself¹⁴⁰

Generative AI models include exact copies of some copyrighted works or portions thereof. Many of the AI Company Comments acknowledge that AI models have a tendency to generate outputs that are nearly identical (in whole or in part) to at least some of their inputs.¹⁴¹ One recent paper

¹³⁵ See, e.g., Adobe Initial Comments at 3-4 (models create mathematical “weights” [not ideas] that help them produce a satisfactory output); Meta Initial Comments at 3-6 (describing process for turning sentences into a matrix of numbers).

¹³⁶ See, e.g., Nasr Article at 14 (“our paper suggests that training data can easily be extracted from the best language models of the past few years through simple techniques”).

¹³⁷ OpenAI Initial Comments at 12; see also, e.g., BSA Initial Comments at 2-3; R Street Initial Comments at 4; Engine Initial Comments at 5; Microsoft Initial Comments at 7-8; Anthropic at 7; Andreessen Horowitz Initial Comments at 7; Authors Alliance Initial Comments at 9, 12-13.

¹³⁸ See Nasr Article at 15 (“In order for GPT-Neo 6B to be able to emit nearly a gigabyte of training data, this information must be stored somewhere in the model weights”); see also Jbetker, *The “it” in AI models is the dataset*, Non-Interactive – Software and ML, June 10, 2023, <https://nonint.com/2023/06/10/the-it-in-ai-models-is-the-dataset/> (“model behavior is not determined by architecture, hyperparameters or optimizer choices. It’s determined by your dataset, nothing else”).

¹³⁹ See, e.g., Center for AI and Digital Policy Initial Comments at 6 (quoting an October 2023 article in Scientific American that found that “[m]achine learning models are only capable of pumping out images and text because they’ve been trained on mountains of real people’s creative work, much of it copyrighted”).

¹⁴⁰ See *Disney v. VidAngel*, *supra* at n.62.

¹⁴¹ See, e.g., Andreessen Horowitz Initial Comments at 6 (acknowledges memorization but claims rates of occurrence are low); Anthropic Initial Comments at 7 (acknowledges that memorization is possible in certain circumstances); BSA Initial Comments at 11 (acknowledges that memorization happens but calls it an “edge case”);

finds that this phenomenon – sometimes anthropomorphized as “memorization,” but more accurately called “copying” – “is more prevalent than previously believed and will likely get worse as models continue to scale, at least without active mitigations.”¹⁴² For example, academic researchers studying Stable Diffusion have found that it is capable of reconstructing some of its training images nearly identically.¹⁴³ Other academic researchers have identified numerous literary works represented in ChatGPT, GPT-4, and other large language models.¹⁴⁴ This is also a basis for the lawsuit recently filed by Concord and other music publishers against Anthropic, in which the plaintiffs allege that the Claude AI system is capable of generating verbatim or near-verbatim copies of song lyrics as outputs even when no effort is made to have it do so.¹⁴⁵ The AI Company Commenters may wish to dismiss this phenomenon as “a bug not a feature,”¹⁴⁶ but that is entirely irrelevant to the question of whether ingested content is reproduced and embodied in the model. The fact that copies of copyrighted works can be reproduced as outputs of an AI model is itself proof that those works were copied and embedded within the model.

Filtering output does not absolve infringers from liability for ingestion. To avoid liability for infringing outputs, some AI developers are developing and implementing techniques to minimize the chances that an AI-generated output will be substantially similar (or identical) to a copyrighted work that it ingested.¹⁴⁷ “These measures have included contractual use restrictions, technical measures like prompt- and output-filters, and warnings that models are not intended to be used to infringe copyright or other intellectual property rights.”¹⁴⁸ Although we salute efforts by AI developers to prevent their systems from generating identical or substantially similar outputs, these efforts do not solve the issue of infringement at the ingestion stage; in fact, these filtering measures have the effect (whether intentionally or unintentionally) of concealing the unauthorized reproduction of copyrighted works that may have happened at the ingestion stage. This result underscores the critical need for two pieces of legislation, both of which are

CCIA Initial Comments at 15-16 (acknowledges memorization occurs but attributes it to inadvertent distortions in the model); EFF Initial Comments at 4-5 (acknowledges memorization but refers to it as a “bug, not a feature”); Google Initial Comments at 13 (“We are aware of concerns that AI systems can output content that is substantially similar to individual pieces of content on which they were trained”); Hugging Face Initial Comments at 7 (describing research into “the tendency of large language models to generate long exact sequences from their training corpus”); Microsoft Initial Comments at 8 (noting that “outputs of an AI model that are substantially similar to their copyrighted work . . . can potentially occur”); OpenAI Initial Comments at 8-11 (includes a whole section describing efforts to prevent memorization or repetition); Samuelson Initial Comments at 9 (acknowledges that “in limited situations, Generative AI models do copy the training data”).

¹⁴² Carlini, *et al.*, *Quantifying Memorization Across Neural Language Models*, arXiv (Mar. 6, 2023), available at <https://arxiv.org/abs/2202.07646>.

¹⁴³ See Nicholas Carlini *et al.*, *Extracting Training Data from Diffusion Models* at 5, arXiv (Jan. 30, 2023), available at: <https://browse.arxiv.org/pdf/2301.13188.pdf>; Gowthami Somepalli *et al.*, arXiv at 1, 8, arXiv (Dec. 12, 2022), available at <https://browse.arxiv.org/pdf/2212.03860.pdf>.

¹⁴⁴ See, e.g., Chang *et al.*, *Speak, Memory: An Archaeology of Books Known to ChatGPT/GPT-4*, arXiv (Oct. 20, 2023), <https://arxiv.org/abs/2305.00118>; Nasr Article.

¹⁴⁵ *Concord Music Grp., Inc. v. Anthropic PBC*, No.3:23-cv-01092 (M.D. Tenn. Oct. 18, 2023), Compl. (ECF No. 1).

¹⁴⁶ See, e.g., EFF Initial Comments at 4-5; OpenAI Initial Comments at 7; BSA Initial Comments at 11, n.32; Google Initial Comments at 13; International Center for Law & Economics Initial Comments at 19.

¹⁴⁷ See, e.g., Chamber of Progress Initial Comments at 10 “For instance, Stability AI has modified its image generator, Stable Diffusion, to deny requests that mimic established artists’ styles. OpenAI made a parallel move with DALL-E 3, which now also refuses to generate images mirroring the style of any living artist and allows artists to opt-out of having their works included in any training sets”).

¹⁴⁸ TechNet Initial Comments at 7; see also, e.g., Anthropic Initial Comments at 7 (output filtering); OpenAI Initial Comments at 7 (training models to decline to respond to certain prompts).

mentioned in the A2IM-RIAA Initial Comments, that will provide copyright owners with access to the data they need to prove infringement at the ingestion stage¹⁴⁹: (1) legislation that imposes detailed recordkeeping and disclosure obligations on AI developers (and others in the AI supply chain)¹⁵⁰ (discussed further below), and (2) legislation that creates a new administrative subpoena, as described in detail in the A2IM-RIAA Initial Comments (including the proposed legislative language included in Annex B, thereto).

IV. Direct and Secondary Liability

AI Company Commenters would make end users solely responsible for infringing outputs. The AI Company Comments acknowledge that AI-generated outputs can (and sometimes do) infringe the exclusive rights in preexisting copyrighted works.¹⁵¹ But rather than accept any responsibility for these infringements, the AI Company Commenters would have all infringement liability lie with end users,¹⁵² apparently even when those users do nothing to deliberately generate an infringing output.¹⁵³ At the same time, many of the AI Company Commenters –

¹⁴⁹ See, e.g., Chamber of Progress Initial Comments at 8 (“rights holders must substantiate their claims by pinpointing the exact infringed works”).

¹⁵⁰ It would also be helpful to have executive action that required AI developers to have such recordkeeping and disclosure obligations with respect to any AI systems licensed or sold to the U.S. government.

¹⁵¹ See, e.g., Anthropic Initial Comments at 11 (“Specific user-generated outputs implicate the copyright in pre-existing works”); Microsoft Initial Comments at 8, 10-11 (“When a person uses an AI application to create expressive works, it is possible that AI generated outputs may infringe copyright if the output is substantially similar to a previous work”); BSA Initial comments at 11 (“the more data available for training, the less likely the system will produce a copy or derivative of an input (in the absence of a user’s intent to infringe)”; EFF Initial Comments at 4-5 (users may “prompt an AI to generate an infringing work”); OpenAI Initial Comments at 8, 10, 14; (“in the rare situations where an output satisfies copyright law’s substantial similarity test, a specific output could implicate the exclusive rights of a copyright owner, depending on the context and facts involved”); Samuelson Initial Comments at 28-29 (“AI-generated outputs implicate the exclusive rights of preexisting copyrighted works when those outputs are substantially similar to specific copyrighted works in the training data”); CCIA Initial Comments at 8 n.9, 20-21; Chamber of Progress Initial Comments at 6, 8, 10-11, 14, 16; Google Initial Comments at 11, 13; Hugging Face Initial Comments at 7, 9;

¹⁵² See, e.g., OpenAI Initial Comments at 14 (“In evaluating claims of infringement relating to outputs, the analysis starts with the user. After all, there is no output without a prompt from a user, and the nature of the output is directly influenced by what was asked for”); Google Initial Comments at 13-14 (“When an AI system is prompted by a user to produce an infringing output, any resulting liability should attach to the user as the party whose volitional conduct proximately cause the infringement”); EFF Initial Comments at 4-5 (“any liability should reside with that user, given that they have done the illicit copying”); CCIA Initial Comments at 8-10, 21 (“any liability should lie on the end-user who requests and publishes a copyright-infringing work”); Anthropic Initial Comments at 12 (“Generally, responsibility for a particular output will rest with the person who entered the prompt to generate it. That is, it is the user who engages in the relevant “volitional conduct” to generate the output and thus will usually be the relevant actor for purposes of assessing direct infringement”) (citation omitted); Chamber of Progress Initial Comments at 6, 11 (“users might craft prompts that sidestep the protective measures set by AI providers, potentially leading to infringement”); see also, TechNet Initial Comments at 6; Microsoft Initial Comments at 10-11; Meta Initial Comments at 9.

¹⁵³ While some users may set out to intentionally generate infringing outputs (e.g., by prompting an AI tool to produce the lyrics to a specific song or the words to a specific poem), despite the AI Company Commenters’ attempts to suggest otherwise, see, e.g., CCIA Initial Comments at 21, EFF Initial Comments at 5, a user does not need to intend to infringe in order to end up with an infringing output. See, e.g., OpenAI Initial Comments at 15 (“the overwhelming majority of users have no interest in infringing any exclusive rights in any pre-existing copyrighted work”). Moreover, users typically have limited (if any) control over the design and behavior of the AI systems they choose to utilize for a particular task.

Adobe, Amazon, Google, IBM, Microsoft and OpenAI to name a few – are offering to indemnify their users against that same liability.¹⁵⁴

The Sony Betamax case does not insulate AI companies against secondary liability. According to AI Company Commenters, end users are the only ones who can be guilty of infringement because AI tools have “substantial non-infringing uses” and the developers and distributors (and others in the supply chain) are, therefore, immune from liability by virtue of the *Sony Betamax* case.¹⁵⁵ This argument misreads *Sony Betamax*, which has no application here.

Contrary to the suggestion of many of the AI Company Commenters, *Sony Betamax* did not hold that distributors of a product are immune from secondary liability so long as the product has substantial non-infringing uses. Rather, as the Supreme Court explained in *Grokster*, “*Sony* barred secondary liability based on presuming or imputing intent to cause infringement solely from the design or distribution of a product capable of substantial lawful use, which the distributor knows is in fact used for infringement” (emphasis added).¹⁵⁶ As the Court made clear, “*Sony* did not displace other theories of secondary liability” and the case “was never meant to foreclose rules of fault-based liability derived from the common law.”¹⁵⁷ To read *Sony* “to mean that whenever a product is capable of substantial lawful use, the producer can never be held contributorily liable for third parties’ infringing use of it,” the Court held, would be “error, converting the case from one about liability resting on imputed intent to one about liability on any theory.”¹⁵⁸

Moreover, the “staple article of commerce” doctrine that *Sony Betamax* applied has no application where, as here, the contributory infringer maintains an ongoing relationship with the party who would be deemed the “direct infringer” under the AI Company Commenters’ approach.¹⁵⁹ Unlike the distributor of a product, AI companies maintain ongoing relationships

¹⁵⁴ See, e.g., Robert Freedman, *Google Is Latest Company To Indemnify AI Users Against Infringement*, Legal Dive (Oct. 13, 2023) [https://techcrunch.com/2023/10/06/some-gen-ai-vendors-say-theyll-defend-customers-from-ip-lawsuits-others-not-so-much/](https://www.legaldive.com/news/google-indemnify-generativeai-users-against-infringement-copyright-ip-genai/696593/#:~:text=Both%20Google%20and%20Microsoft%20are,its%20AI%20assisted%20design%20tools; Kyle Wiggers, <i>Some Gen AI Vendors Say They’ll Defend Customers From IP Lawsuits. Others, Not So Much</i>, TechCrunch, (Oct. 6, 2023 11:21 AM) <a href=); *New Models and Developer Products Announced at Dev Day*, OpenAI Blog dated Nov. 6, 2023) (OpenAI announces its new “copyright shield” saying that “we will now step in and defend our customers, and pay the costs incurred, if you face legal claims around copyright infringement”), <https://openai.com/blog/new-models-and-developer-products-announced-at-devday>. See also, Microsoft Initial Comments at 11. Indeed, such indemnification offers might in certain circumstances be considered to amount to a form of inducement by the AI developers for users to use their AI tools to engage in acts of infringement.

¹⁵⁵ *Sony Corp. of Am. v. Universal City Studios, Inc.*, 464 U.S. 417 (1984). See, e.g., Samuelson Initial Comments at 31 (“Courts may need to consider whether Generative AI systems are technologies that have substantial non-infringing uses and hence eligible for the Sony safe harbor that the Court reaffirmed in *MGM v. Grokster*”) (citations omitted).

¹⁵⁶ *Metro-Goldwyn-Mayer Studios Inc. v. Grokster, Ltd.*, 545 U.S. 913, 933 (2005).

¹⁵⁷ *Id.* at 934-35 (citation omitted).

¹⁵⁸ *Id.* at 934.

¹⁵⁹ See *Cartoon Network LP, LLLP v. CSC Holdings, Inc.*, 536 F.3d 121, 133 (2d Cir. 2008) (recognizing that lack of “ongoing relationship” between product distributor and product user was critical factor in *Sony*); *BMG Rights Mgmt v. Cox Commc’ns, Inc.*, 199 F. Supp. 3d 958, 973-76 (E.D. Va. 2016) (“An ongoing relationship between a

with the users of their generative AI models (in addition to the models themselves) and thus are far from the limited holding of *Sony Betamax*.

Simply put, regardless of whether generative AI systems are capable of substantial non-infringing uses, normal principles of direct and secondary liability will and should appropriately place the responsibility for infringement with AI developers (and others in the supply chain) and/or users. Such determinations are fact-dependent and, as in other contexts, liability will be assessed based on the evidence adduced regarding who actually engaged in infringing acts and who, if anyone, meets the relevant thresholds for secondary liability, including vicarious, contributory, or inducement liability.¹⁶⁰ Nothing in *Sony Betamax* is to the contrary.

Existing safe harbors do not apply and no new safe harbor should be created. In another attempt to escape liability for their actions, some AI Company Commenters suggest that the Office recommend a new safe harbor be created to insulate them from secondary liability.¹⁶¹ In making this suggestion, the AI Company Commenters acknowledge that existing safe harbors, such as Section 512, do not apply to generative AI systems.¹⁶² The Kernochan Center cogently explains why this is so:

The systems are not host service providers ‘stor[ing] at the direction of a user [infringing] material that resides on a system or network controlled . . . by or for the service provider.’ The training data is not user-posted content, and, while end-users enter prompts, the prompts themselves are not infringing content stored at the behest of the user. The infringing content may be incorporated in the outputs, but the system itself generates the outputs. As a result, common law principles of secondary liability would apply, and can evolve to fit the circumstances.¹⁶³

The Office is well aware of the many problems with the existing Section 512 safe harbor.¹⁶⁴ Among other things, where it applies, Section 512 has made it extremely difficult for rights owners to enforce their rights. This is especially true for smaller rights owners without the resources to implement the type of programs necessary to implement an effective notice-and-takedown program. Given how poorly Section 512 has worked, there is no reason to consider expanding the existing safe harbors or creating a new AI-specific safe harbor. We would staunchly oppose any such recommendation.

defendant and direct infringers presents a potential for culpability quite beyond” what *Sony Betamax* represents); *Arista Records LLC v. Usenet.com, Inc.*, 633 F. Supp. 2d 124, 156 (S.D.N.Y. 2009) (“Defendants maintain an ongoing relationship with their users. . . . *Sony’s* insulation from contributory liability is inapplicable”).

¹⁶⁰ See, e.g., *A&M Records, Inc. v. Napster, Inc.*, 239 F.3d 1004, 1019, 1022 (9th Cir. 2001); *Metro-Goldwyn-Mayer Studios Inc. v. Grokster, Ltd.*, 545 U.S. 913, 936 (2005); see also 3 Melvin Nimmer & David Nimmer, Nimmer on Copyright § 12.04 (2020).

¹⁶¹ See, e.g., Chamber of Progress Initial Comments at 12.

¹⁶² *Id.* at 12-13 (“Section 230 is not available for Copyright infringement claims”); Katherine Lee *et al.* Initial Comments at 91 (p. 90 of attached paper) (“Section 512 of the Copyright Act, enacted as part of the Digital Millennium Copyright Act, overlays safe harbors for certain online intermediaries on to copyright law. Although these safe harbors have been significant for technology platforms and for Internet law, none of them are likely to apply to generative AI in most cases”) (citations omitted).

¹⁶³ Kernochan Center Initial Comments at 13-14 (citations omitted).

¹⁶⁴ See U.S. Copyright Office, Section 512 of Title 17 (2020), <https://www.copyright.gov/policy/section512/section512-full-report.pdf>.

V. Transparency and Recordkeeping

Copyright owners need access to adequate records regarding the works ingested by AI models.

As stated in the A2IM-RIAA Initial Comments and in the comments the Creative Community filed with OSTP¹⁶⁵ and NTIA,¹⁶⁶ it is essential that all parties involved in the AI supply chain (including, but not limited to, AI deployers, AI developers, and those who collect and curate copyrighted works for ingestion by AI tools) be required to collect, retain, and disclose records regarding the materials ingested by their models and how those materials are used.¹⁶⁷ Without such records, copyright owners will often be unable to enforce their rights¹⁶⁸ and they will lack the baseline data needed to develop royalty structures and other key terms needed for voluntary licensing. In fact, a variety of companies and groups within the tech sector seem to regard transparency as broadly beneficial and achievable.¹⁶⁹

Copyright Office support for recordkeeping and disclosure is critical. Copyright Office support for recordkeeping and disclosure rules is essential not just because such rules are demonstrably necessary, but because the perceived lack of such obligations in the U.S. is beginning to be used as an argument against recordkeeping obligations in other countries.¹⁷⁰ These recordkeeping and disclosure mandates could be implemented to some extent via Executive Order¹⁷¹ – and that might be the most expedient first step. In addition, federal legislation will likely be needed to secure comprehensive requirements and effective enforcement mechanisms.

The arguments against recordkeeping are easily refuted. To the extent that some AI Company Commenters oppose recordkeeping, their arguments are largely self-serving and unavailing.

¹⁶⁵ See Annex C to A2IM-RIAA Initial Comments.

¹⁶⁶ See June 12, 2023 Comments of the Creative Community in Response to NTIA RFC on AI Accountability (“Creative Community NTIA Comments”) available at <https://www.regulations.gov/comment/NTIA-2023-0005-1277>.

¹⁶⁷ Although not a recordkeeping provision itself, the administrative subpoena proposed in the A2IM-RIAA Initial Comments, see pages 11, 31 and Annex B, is related to the lack of transparency that rights owners face when trying to enforce their rights.

¹⁶⁸ See, e.g., Center for AI and Digital Policy Initial Comments at 11 (“without extensive documentation and recordkeeping requirements it will be impossible to assess whether or not copyrighted material has been used in training data, to what extent it has been used, and any potential infringement of existing copyrights by AI-generated outputs”); see also discussion on pages 22-23, *supra*, regarding the heightened need for recordkeeping and disclosure when AI developers employ prompt and output filters to lessen the likelihood of substantially similar outputs.

¹⁶⁹ See, e.g., Hugging Face Initial Comments at 2,12; Public Knowledge Initial Comments at 15; see also Comments Filed by Google in Response to National Telecommunications and Information Administration AI Accountability Policy Request for Comment (“NTIA RFC on AI Accountability”) at 23 <https://www.regulations.gov/comment/NTIA-2023-0005-1308> (“AI developers can support accountability by retaining detailed documentation of datasets and Models”); Microsoft’s Comments in Response to NTIA RFC on AI Accountability at 1, 3 <https://www.regulations.gov/comment/NTIA-2023-0005-1337> (“Transparency around AI systems is critical to society being able to understand and scrutinize the impact of AI and to ensuring that those using AI do so in a responsible and accountable manner;” also noting that transparency is one of Microsoft’s six responsible AI principles).

¹⁷⁰ For example, the absence of U.S. recordkeeping standards is currently being cited as a reason to remove such obligations from the draft E.U. AI Act.

¹⁷¹ Executive Order, Section 5.2(c)(iii), at 75207.

Neither the volume of the data¹⁷² nor the allegedly competitively sensitive nature of the works used for ingestion are reasons to avoid recordkeeping and disclosure obligations in appropriate circumstances.¹⁷³ With respect to the amount of copyrighted works at issue, the task of recordkeeping can be easily automated, especially if AI developers are careful not to remove the copyright management information (“CMI”) or other metadata that typically accompanies digital copies of copyrighted works. And concerns about competitive sensitivity can be addressed by limiting access to sensitive information to stakeholders with a legitimate need to know (e.g., regulators, copyright owners). The approach we recommend helps copyright owners enforce their rights against existing, widespread infringement while providing guardrails for speculative concerns about potential trade secrets.

VI. Licensing

Voluntary licensing of sound recordings is common, widespread, and demonstrably feasible in the context of AI. Many AI Company Commenters argue that it would be impracticable to clear rights to the large quantities of copyrighted works that AI developers scrape from the Internet,¹⁷⁴ but this is not the way legitimate uses of copyrighted works normally work. The law requires users to license works first, and only then acquire (lawful) copies of them. These arguments against voluntary licensing are particularly flawed when applied to sound recordings.

- First – and most important – as discussed in greater detail in the A2IM-RIAA Initial Comments, voluntary licensing of tens of millions of individual sound recordings owned by thousands of different rights owners is commonplace. After decades of experience, the entire recorded music sector, major and independent labels as well as a vast number of individual rightsholders and independent creators, has the necessary infrastructure to support robust voluntary licensing of a huge range of businesses and business models – from sampling to sync licenses to streaming to social media and fitness apps. In fact, a number of the AI Company Commenters already have fully licensed music streaming services and/or platforms that include music, which provides them with the necessary relationships and know-how to facilitate licensing discussions regarding their AI services. In addition, a market for licensing recorded music for AI uses is already emerging.¹⁷⁵
- Second, the quantity of music needed for an AI model is significantly less than that needed for textual large language models, and recent experience suggests that a smaller

¹⁷² See, e.g., Google Initial Comments at 11-12.

¹⁷³ *Id.*; Adobe Initial Comments at 5; TechNet Initial Comments at 12.

¹⁷⁴ See, e.g., Andreessen Horowitz Initial Comments at 10 (“[Many] AI models . . . were trained on an enormous cross-section of all of the publicly available information ever published on the internet—that is, billions of pieces of text from millions of individual websites. For a very significant portion of those works, it is essentially impossible to identify who the relevant rights holders are, and thus there would be no viable way to get statutory royalties to the proper parties”); Anthropic Initial Comments at 10 (most works scraped from the Web do not have management information); CCIA Initial Comments at 15 (many of the works ingested by AI “lack any identified or identifiable author from whom to obtain a license”); EFF Initial Comments at 4 (authorization from every copyright owner not feasible due to copyright attaching to creations at time of fixation).

¹⁷⁵ See, e.g., pages 1-2 *supra*; see also A2IM-RIAA Initial Comments at 13, n.45.

number of quality inputs will yield higher quality AI outputs than an AI system that ingested a wider swath of works of varying quality.¹⁷⁶

- Third, it is up to copyright owners, not AI companies, to decide whether a royalty rate is sufficient to justify a deal and to structure their license agreements in ways that make economic sense for them. The fact that the AI Company Commenters think the rate would be too low for licensors provides no basis for sheltering would-be licensees from their obligation to license copyrighted works before exploiting them.¹⁷⁷

There is no support for compulsory licensing or ECL. Finally, as noted in our introduction, there is virtually no support from would-be licensees for a compulsory license (or for extended collective licensing (“ECL”). Indeed, the AI Company Comments generally oppose statutory licensing as well.¹⁷⁸ Given the broad opposition of both the tech sector and copyright owner groups,¹⁷⁹ we urge the Office to definitively reject this idea in its forthcoming report.

Conclusion

Once again, we thank the Office for raising so many important and timely issues related to the intersection of copyright law and generative AI. We also appreciate the Office’s willingness to grant a one-week extension to the deadline for these Reply Comments. As the Office undertakes its review of the nearly 10,000 comments in the initial round and the thousands more that will be filed in this reply round, we encourage it to embrace as its guiding principle the value of human creativity and the important role that copyright law has always played in incentivizing creators to pursue their arts.

December 6, 2023

¹⁷⁶ See, e.g., Matt Mullen, *AI music wars: Meta takes on Google and releases its own AI music generator – but whose is better?*, MusicRadar (June 16, 2023), <https://www.musicradar.com/news/meta-google-ai-music-wars-musicgen> (last visited Oct. 23, 2023) (concluding that Meta’s product, which is trained on significantly less music than Google’s product, creates better music); Jade Copet *et al.*, *Simple and Controllable Music Generation*, arXiv, June 8, 2023, available at <https://arxiv.org/pdf/2306.05284.pdf>.

¹⁷⁷ See, e.g., Anthropic at 10 (“it will be difficult to set a royalty rate that is meaningful to individual creators without making it uneconomical to develop generative AI models in the first place”); Meta Initial Comments at 20 (“any fair royalty would be incredibly small in light of the insignificance of any one work among an AI training set”); Hugging Face Initial Comments at 11,

¹⁷⁸ See, e.g., Andreesen Horowitz Initial Comments at 8; CCIA Initial Comments at 15; TechNet Initial Comments at 9; Authors Alliance Initial Comments at 1, 14-16; Duolingo Initial Comments at 3; Engine Initial Comments at 8; Peermusic and Boomy Initial Comments at 6; Software Freedom Conservancy Initial Comments at 6; Microsoft Initial Comments at 9; R Street Initial Comments at 6; Creative Commons Initial Comments at 6; Consumer Technology Association Initial Comments at 5; Meta Initial Comments at 20.

¹⁷⁹ See, e.g., Copyright Alliance Initial Comments at 80; MPA Initial Comments at 28-29; NMPA Initial Comments at 24; AAP Initial Comments at 25; News Media Alliance Initial Comments at 53-54; Copyright Clearance Center Initial Comments at 13; Recording Academy Initial Comments at 8; ASCAP Initial Comments at 40-43; Universal Music Group Initial Comments at 66-67.